

PLAYING HARD TO GET: STRATEGIC SIGNALING IN AID BARGAINING*

Anthony Luongo[†]

November 14, 2025

Abstract

Foreign aid is a political exchange between a donor and target. Existing literature focuses primarily on the strategic actions of donors, but less is known about how targets advance their interests. This article explores the role of aid-receiving states as strategic actors who can extract concessions from donors. I model the aid exchange using a costly signaling model in which targets send a (potentially misleading) signal of their policy preferences before the donor makes an aid offer. In equilibrium, when the cost of sending a misleading signal is sufficiently low, targets who are aligned with a donor on policy sometimes lie about their alignment, which yields them aid that they would not have received otherwise. Empirically, I show that strategic nonresponse in the UN General Assembly – a low-cost signal of nonalignment – is correlated with higher future aid inflows.

Word count: 11,998

*I thank Sam England, Jennifer Gandhi, Nahomi Ichino, Stefano Jud, Jack Paine, John Patty, Eric Reinhardt, Miguel Rueda, Renard Sexton, and Jessica Sun for their advice and feedback on this project. I am also grateful to participants in the Emory International Relations Journal Club (2023), Southeastern Rights Network Meeting (Winter 2023), the Midwest Political Science Association Annual Conference (2024), the Political Economy of International Organization Annual Conference (2024), and the Political Economy of Aid Society Meeting (2024) for their helpful suggestions.

[†]Ph.D. Candidate, Department of Political Science, Emory University

The 2022 Russian invasion of Ukraine sparked swift backlash from Western states, who placed harsh sanctions on the Russian economy and provided material assistance to Ukraine's defense efforts. However, not all Western partners followed this lead; Vietnam and South Africa – not necessarily formal allies, but states that extensively cooperate on economic and security issues with the West – joined US adversaries like China and Iran in abstaining from a UN General Assembly (UNGA) vote that condemned the invasion, although they did not endorse Russia's position (Voeten, 2022). Beyond the UNGA, South Africa hosted Russian Foreign Minister Sergei Lavrov for a visit and participated in joint military exercises with Russia, even as it continued cooperating with its Western partners as normal (Allison, 2023; Eligon, 2023; Makhanya, 2023). Despite Western diplomats' claims that they were not concerned about where South Africa's true loyalties lie, the public friendliness with Russia appears to have sparked concern, with the US quickly arranging several high-profile diplomatic visits (Gramer, 2023).

Why would South Africa and other states with deep Western ties abstain on this vote? Material considerations such as oil and arms likely played a role. However, it seems unlikely that these states view Russia as a viable long-term partner; Russia is an increasingly isolated state with a precarious economy and insignificant development assistance program, and South Africa does not lack other options. It also seems unlikely that South Africa – led by the African National Congress, a party with a history of anticolonial activism – views Russia's imperial project as genuinely righteous.

In this article, I argue that aid-receiving states like South Africa can exaggerate the distance between their policy ideal point and the policy proposal of a donor to incentivize the donor to give them a level of aid that exceeds their true reservation price. This argument posits a novel explanation for the policy divergence between the US and some of its partners on the issue of Russia, framing UNGA abstentions and Russian diplomatic visits as an attempt by those partner states to attract more resources from the US. The flurry of diplomatic activity on the US side is early evidence that

this strategy is bearing fruit for South Africa.

I focus on this strategy in foreign aid bargaining, a form of political exchange between donors and recipient states. Donors use aid to stabilize allied regimes, incentivize domestic policy changes by partner states, and purchase support for diplomatic initiatives, among other things (Alesina and Dollar, 2000; Dreher, Lang, et al., 2022). Recipients use aid to enhance their political position among elites and/or key voting constituencies (Cruz and Schneider, 2017; Morrison, 2009).

In general, recipients want to maximize aid inflows while minimizing policy concessions; donors want to achieve certain policy goals while minimizing aid spending (Bueno de Mesquita and Smith, 2016). However, donors may not know the target's true policy preferences and thus cannot discern the amount of aid necessary to ensure its preferred policy is implemented. In this article, I focus on how target states exploit the donor's dilemma through strategic signaling behavior. This contrasts with most existing work on aid allocation, which focuses on the role of donor interests and structural features of recipients. I argue that target states that are aligned with the donor may portray themselves as *less* aligned with the donor to receive more aid. This works by inflating the donor's perception of the target's reservation price for implementing the donor's preferred policy. Consequently, the donor believes – inaccurately – that the target will not do what the donor wants without aid as an incentive.

I develop my argument using a costly signaling model before testing it empirically. In the model, the donor does not know whether the recipient shares their policy preference. The recipient reveals information about this preference before the donor makes an offer of aid; this signal is costly if and only if the target lies about their preference. Analysis of the model shows that, when the cost of lying is sufficiently low and donor's interest in the target's policy is sufficiently high, aligned targets optimally claim nonalignment and donors, uncertain over the target's true preferences, give those states aid at least some of the time. By contrast, with complete information, donors would recognize that such aid is unnecessary for policy alignment. I derive and test implications of the model using

data on nonresponse (abstentions and absences) in the UN General Assembly (UNGA), a forum in which signals can be sent clearly and at relatively low cost. I argue that UNGA resolutions are a tool for donors to infer partners' underlying preferences and likely future actions; because UNGA votes are often inconsequential in their own right, recipients do not face high costs to lying in those votes. Using a selection on observables approach, I show that strategic nonresponse is significantly and positively associated with future aid from the US and US-dominated international institutions for poor US allies. The increase is substantively meaningful; a one standard deviation increase in strategic nonresponse is correlated with about a \$71 million aid increase. Consistent with the model, the returns to nonresponse decline as foreign policy similarity increases.

Given considerable power imbalances between donors and recipient and the asymmetric nature of the aid relationship, existing scholarship on who gives aid to whom has focused on donors' interests as the drivers of aid allocation (Alesina and Dollar, 2000; Dang and Stone, 2021; Scott and Carter, 2019; Steinwand, 2015). However, there is a growing recognition that recipients can influence the amount and terms of aid in their favor under particular circumstances, including when they have large economies (Bayer et al., 2015), hold key positions in international institutions (Berlin et al., 2022), present spillover risks to the donor's economy (Ferry and Zeitz, 2024), or have access to alternative sources of revenue (Brown, 2023; Horning, 2008; Winters, 2024). More broadly, this work fits within a growing interest in the "power of the weak", or how geopolitically minor states advance their interests (Snidal et al., 2024). This article contributes to this growing literature by proposing a novel theoretical pathway – the misrepresentation of their preferences – through which recipients can improve aid bargaining outcomes and empirically testing the implications of that theory. To my knowledge, this article is the first to explicitly model aid recipients as actors who can behave strategically beyond the decision to accept or reject aid.

Misrepresentation of policy preferences is a type of signaling action, a class of behavior that

has received considerable attention in the broader IR literature, but has not yet been explored in the aid context. In signaling setups, a player *A* has a type that they observe, but the other player *B* does not observe; the realization of this type affects the utility of both players and thus the actions they choose to take. *A* sends a signal to *B* that reveals information about that type, but might not be honest; *B* updates its beliefs about the type accordingly. In crisis bargaining, state *A* wishes to signal that they are resolved (their type) by taking costly actions like mobilizing their military, which might lead *B* to believe that *A* is indeed resolved and thus grant *A* concessions (Wolford, 2014).¹ Despite the prevalence of signaling models and the argument of major IR theorists that incomplete information is a fundamental feature of international politics (e.g., Fearon, 1995), existing work on aid bargaining uniformly assumes that donors know recipients' policy preferences and thus how much aid is necessary to gain concessions (e.g., Bueno de Mesquita and Smith, 2016). The logic of preference misrepresentation is also consistent with case evidence. For example, Uganda held a seat on the UNSC when the major powers sought to enact new sanctions against Iran over its nuclear weapons program. Just three weeks before a vote on the sanctions, Uganda hosted Iranian president Ahmadinejad for a state visit to discuss the issue. Uganda did eventually vote in favor of the sanctions, but it was granted concessions on crucial security issues by the Security Council (Mikulaschek, 2021). It seems unlikely that Uganda had any true interest in protecting Iran's nuclear program, instead giving that impression to attract concessions.

In the empirical section of this paper, I argue that the UNGA is a suitable forum for preference misrepresentation, suggesting a novel informational function of the institution. Existing work generally follows Dixon (1981) in painting the UNGA as "a passive arena for the political interaction of member states" (p. 47). This view has led scholars to view the institution as generally inconsequential for international politics in its own right and as an imperfect but suitable forum

¹Signaling games are also commonly analyzed in the literature on economic sanctions, escalation in interstate disputes, and terrorism (Crisman-Cox and Gibilisco, 2021).

for estimating state's underlying preferences (Bailey et al., 2017; Voeten, 2000). My argument is consistent with the former point but inconsistent with the latter because I show it is a forum for misrepresentation of preferences. I also diverge from the work on aid that assumes that donors are interested in buying UNGA votes with aid (Dreher, Nunnenkamp, and Thiele, 2008). Instead, I argue that the limited policy effects of the UNGA lend the forum to a different purpose for developing states in particular: it serves as a forum for signaling policy preferences – not necessarily truthfully – in order to attract resources from donor states. Donors use the UNGA as a means of inferring a target's preferences over a certain policy, which are distinct from but related to the content of an UNGA vote. This raises questions about the validity of UNGA-derived measures of alignment, which are common in studies of international politics (e.g., Tomashevskiy, 2021). Narrowly, the model suggests that UNGA voting underestimates the degree of alignment between aid-receiving states and their leading donor(s). More broadly, it points to the possibility that UNGA votes are used strategically to pursue goals unrelated to the resolution in question.

Theoretical Model

I build a costly signaling model of aid-for-policy deals.² Existing work on aid-for-policy deals comes from the selectorate literature and focuses largely on how structural features of states shape foreign aid (Bueno de Mesquita and Smith, 2016). My model builds on this work by relaxing the assumption that donors know the target's policy preferences. Instead, donors have incomplete information over the target's ideal point; donors have beliefs over the probability that the target is aligned with them, but do not know the realization of that probability. Targets exploit this lack of information by misrepresenting their type to convince donors that they require aid to implement a certain policy, even though receiving aid would not increase their likelihood of implementing that

²Broadly speaking, aid has two political functions: purchasing policy and subsidizing favorable regimes. These two functions have different logics and thus it is difficult to model them together. I focus on the first goal and acknowledge that the model fails to predict behaviors that stem from the latter.

policy.

In the model, a donor (D) and a target (T) bargain over a transfer of resources from donor to target (r_D) and a policy to be implemented by T (y_T). The target has some genuine preference (θ_T , T's type in signaling terms) that is either aligned with the donor's preferred policy or unaligned. The donor knows the probability that the target is aligned with them, but does not observe the alignment. Both players care to some degree about the target implementing their most preferred policy; α_D represents the benefit gained by D if T implements her most preferred policy, while α_T is the cost incurred by T implementing a policy other than its most preferred policy. Prior to the offer of aid and policy implementation, the target has an opportunity to send a signal of their view of the donor's policy (s_T); empirically, this could be many things, such as releasing a statement or report of the government's assessment of that policy, votes on non-binding or binding resolutions, or diplomatic visits with proponents or opponents of the policy. This signal does not need to be honest, but the target pays a cost of lying (c); this cost is imposed by the domestic winning coalition, which may interpret the signal as evidence that their own preferred policy is less likely to be implemented. The donor observes the signal – but crucially, does not observe any costs of signaling incurred by the target – and then decides on an allocation of resources to send to the target, conditional on the target's implementation of the donor's preferred policy. After seeing the offer, the target decides which policy to implement. If the donor's preferred policy is implemented, the proposed resource transfer occurs.³

A key assumption is that misrepresentation of type (i.e. sending a dishonest signal) is costly to the target. An assumption that different signals impose distinct costs on the sender is a standard feature of signaling games, and existing models of electoral competition have introduced costs specifically for lying (Callander and Wilkie, 2007). I conceptualize this cost as stemming from

³A rich literature has highlighted commitment problems for both donors and recipients of aid (Curtice and Reinhardt, 2023; Swedlund, 2017). I set that issue aside to focus on information problems.

domestic political consequences of sending a dishonest signal, which allows the realized value to vary by both method of delivery and issue area. Consider the following two-level game illustration of this cost: a political leader acts as an agent for two principals with sometimes-conflicting interests, her aid donor and her winning coalition. Both principals are interested in a signal of the policy she is likely to implement in the future. The donor uses this signal to decide the optimal aid offer, while the winning coalition uses it to decide whether to mobilize against the leader. The extent of this mobilization varies based on the bindingness of the signal and salience of the policy issue. For example, diplomatic cables show that a Tuvaluan diplomat stated his country's goal at the UNGA quite plainly to US diplomat: "We are here to seek assistance" (Plaisted, 2009). By contrast, when Greek Prime Minister George Papandreou suggested to European leaders that his country would hold a referendum on an unpopular, austerity-laden bailout deal – which would in effect mean a referendum on eurozone membership – the backlash from within his government was so severe that Papandreou was forced to step aside (Spiegel, 2014).⁴ In the Tuvaluan case, signaling dishonestly was inconsequential, but for Papandreou's government it was devastating. These divergent anecdotes reflect that lying costs can vary across contexts. Most UNGA issues are likely not salient in Tuvaluan politics, but clearly eurozone membership is quite salient in Greece; lying costs should be higher when an issue is more salient. Costs are also likely higher in democracies, where winning coalition members cannot plausibly be informed about or as clearly compensated for a dishonest signal.

Finally, what policy linkages are donors uncertain about and how do they relate to the target's winning coalition? The answer to this naturally varies across dyad-years. Donors may care about gaining meaningful support on key diplomatic initiatives, access to markets and natural resources

⁴This cost is not observed by the donor. The empirical applicability of this assumption, and thus the broader model, are related to intraregime coordination (i.e. communication that a dishonest signal is aid-seeking), which may stem from regime type and winning coalition size. The unobservability assumption is often reasonable because the model describes routine interactions between states across many issue areas, where a donor deals with multiple states. This is distinct from the crisis bargaining literature (e.g., Ramsay, 2004) where similar costs are observable. This literature describes high-profile interactions between a pair of states in crisis, where expending some effort to discern domestic political dynamics is more often optimal.

for firms from the donor's country, economic and political liberalization, and other institutional reforms like corruption reduction. Members of the target's winning coalition have preferences over these issues. Ideological and/or material ties may shape preferences over international relations; for example, influential business elites with ties to China may pressure leaders not to engage diplomatically with Taiwan. Changes to economic policy also threaten business elites as well as laborers. In the past, Western aid has been conditional on reduction in protectionist policies, exposing domestic firms who had benefited from government protection to competition with more efficient foreign firms (Dang and Stone, 2021). Finally, autocratic elites are likely to oppose political liberalization that upends the political status quo, especially if such liberalization may push them out of power and expose them to retaliation. UNGA resolutions relate to these issues; for example, past UNGA resolutions have condemned both US allies (i.e., Israel) and adversaries (i.e., North Korea), proposed norms and rights for domestic governance, and taken positions on international security issues like arms control and nuclear disarmament. These resolutions need not be perfect mappings of the issue; they just need to convey information about the likely stance of the target.

Model Setup

I now introduce the game more formally. There are two players, $i \in \{T, D\}$, who have the following utility functions:

$$U_T(s_T, y_T, r_D; \theta_T) = r_D \cdot y_T - \mathbb{1}_{y_T \neq \theta_T} \cdot \alpha_T - \mathbb{1}_{s_T \neq \theta_T} \cdot c$$

$$U_D(r_D, y_T) = \alpha_D \cdot y_T - r_D \cdot y_T$$

where $r_D \geq 0$ is the amount of resources transferred from D to R, $y_T \in \{0, 1\}$ is a policy, $\alpha_i > 0$ is the degree of player i 's interest in T's policy, $\theta_T \in \{0, 1\}$ is T's type and corresponds to T's most preferred policy, $s_T \in \{0, 1\}$ is a signal of T's type chosen by T to indicate (not necessarily

honestly) their most preferred policy, and $c > 0$ is the cost associated with sending a dishonest signal ($s_T \neq \theta_T$). I assume that D prefers that $y_T = 1$. The probability that T is aligned with D (i.e. the probability that $\theta_T = 1$) is $p \in (0, 1)$. The game proceeds as follows (which is represented in game tree form in Figure 1):⁵

1. Nature chooses whether the Target is aligned with the Donor's policy preferences ($\theta_T \in \{0, 1\}$).
Nature chooses an aligned type with probability $p = \Pr[\theta_T = 1]$. This alignment is observed only by the Target.
2. The Target privately observes its preference θ_T and selects a signal $s_T \in \{0, 1\}$.
3. The Donor observes the signal (but not θ_T or any cost incurred by the Target) and decides how much aid r_D (if any) to offer the Target, conditional on the Target's implementation of the Donor's preferred policy $y_T = 1$.
4. T observes the aid offer r_D and chooses to implement a policy $y_T \in \{0, 1\}$. If the Target chooses the Donor's preferred policy $y_T = 1$, the Donor pays the promised amount of aid r_D to the Target.
5. Payoffs are realized.

Given this sequence of play, each player has a strategy that maps features of the game and/or other player's actions into their own actions. The target T has a strategy composed of two functions. The first, $\sigma_T^s : \{0, 1\} \rightarrow [0, 1]$ maps a type θ_T into a probability of sending a signal $s_T = 1$. The second, $\sigma_T^y(r_D) : \{0, 1\} \times \mathbb{R}^+ \rightarrow [0, 1]$, maps the type and D's aid offer (i.e. each pair (θ_T, r_D)) into a probability of implementing the donor's preferred policy $y_T = 1$. I refer to these two strategies together as σ_T (that is, $\sigma_T = \{\sigma_T^s, \sigma_T^y\}$). Likewise, a donor's strategy is a function $\sigma_D(s_T) : \{0, 1\} \rightarrow \mathbb{R}^+$ that maps a signal s_T into a weakly positive offer of aid r_D (made before

⁵Payoffs are displayed with Target utility on top and Donor utility on the bottom.




Figure 1: Game Tree

observation of the policy y_T) to be sent if and only if $y_T = 1$. Additionally, β_D denotes D's beliefs over the type θ_T of the target $\beta_D : (0, 1) \times \{0, 1\} \rightarrow [0, 1]$, which defines the posterior probability that $\theta_T = 1$ for each pair (p, s_T) , the prior probability of $\theta_T = 1$ and the signal s_T .

I use perfect Bayesian equilibrium (PBE) as my solution concept. In PBE, an equilibrium is defined by strategies and beliefs that are consistent with Bayes' rule. A strategy profile $(\sigma_T, \sigma_D, \beta_D)$ constitutes an equilibrium whenever σ_T and σ_D are optimal, given the beliefs β_D . I denote an equilibrium strategy profile as $(\sigma_T^*, \sigma_D^*, \beta_D^*)$.

Equilibrium Analysis

The Target will only accept an offer of aid that makes them weakly better off. Otherwise, they will implement their most preferred policy. This does not mean that the Target who does not accept an offer of aid always implements the Donor's least preferred policy; if the Target truly wants to implement $y_T = 1$, then they will regardless of the offer of aid. Lemma 1 presents this strategy more formally. Proofs of this and all formal results are in Appendix B.

Lemma 1 *A Target who receives an offer of aid that is weakly greater than α_T implements $y_T = 1$. A Target who does not receive such an offer implements their most preferred policy $y_T = \theta_T$. The optimal policy implementation strategy σ_T^{y*} is thus*

$$\sigma_T^{y*}(\theta_T, r_D) = \begin{cases} \theta_T & \text{if } r_D < \alpha_T \\ 1 & \text{if } r_D \geq \alpha_T \end{cases}$$

In any equilibrium, there are only two possible offered quantities of aid: zero and α_T . The Target's interest in implementing its most preferred policy (α_T) is also the Target's reservation price for implementing a policy other than its most preferred policy. Because there are only two possible types, there are only two possible reservation prices. An offer between zero and α_T does not convince unaligned types, while an offer that exceeds α_T results in the same outcome as offering α_T . This fact allows the bargaining element of the game to be collapsed from a continuum of actions into only two feasible actions. This also indicates an alternative interpretation of α_T as the Target's demand

for aid.

Lemma 2 *In any equilibrium, the Donor's aid offer equals either 0 or α_T .*

Lemmas 1 and 2 describe behavior that applies to any equilibrium. I now introduce the features of equilibria that vary based on parameter values. This behavior depends on the cost of lying (c), the prior probability of alignment (p), and the relative values of the weight placed on the Target's policy by both the Target and Donor (α_T and α_D , respectively). I introduce the behavior formally in Propositions 1-4, while Figure 2 displays graphically where in the parameter space each equilibrium occurs when the donor is willing to make a strictly positive offer of aid to targets it believes may be unaligned (i.e. $\alpha_D > \alpha_T$). I offer the necessary belief structure for the donor, the optimal signal sent by the target, the optimal aid offer by the donor, and finally the optimal policy to be implemented by the target. I make the simplifying assumption that D is only willing to pay for policy change when it makes them strictly better off (when $\alpha_D > \alpha_T$).

Proposition 1 (Separating Equilibrium) *Suppose that the cost of lying is high ($c \geq \alpha_T$) and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). There is a separating equilibrium in which the Target signals honestly and the Donor pays each type of Target its reservation price. The policy y_T is implemented according to Lemma 1; in equilibrium, both types of Target implement $y_T = 1$.*

$$\begin{aligned}\beta_D^*(p, s_T) &= \begin{cases} 0 & \text{if } s_T = 0 \\ 1 & \text{if } s_T = 1 \end{cases} \\ \sigma_T^{s*}(\theta_T) &= \theta_T \\ \sigma_D^*(s_T) &= \begin{cases} 0 & \text{if } s_T = 0 \\ \alpha_T & \text{if } s_T = 1 \end{cases}\end{aligned}$$

Proposition 1 describes a separating equilibrium, by which I mean each type of Target takes a unique signaling action (in this case sending the signal that matches their type). In this equilibrium, lying is highly costly to the Target, so even though the Donor will respond to an unaligned signal with aid, lying is ultimately not worth the increased aid inflow. Intuitively, this suggests that the Target will not use high stakes policy decisions as opportunities to misrepresent their preferences.

This is the Donor's most preferred equilibrium, as they will give aid equivalent to the Target's true reservation price for all types of the target.

Proposition 2 (Pooling Equilibrium) *Suppose that the cost of lying is low ($c < \alpha_T$), the prior probability of the Target being aligned with the Donor is low ($p \leq 1 - \frac{\alpha_T}{\alpha_D}$), and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). There is a pooling equilibrium in which the Target always signals that they are unaligned with the Donor and the Donor pays all types of the Target the unaligned type's reservation price (α_T). Consistent with Lemma 1, both types of Target implement $y_T = 1$ in equilibrium.*

$$\beta_D^*(p, s_T) = \begin{cases} p \leq 1 - \frac{\alpha_T}{\alpha_D} & \text{if } s_T = 0 \\ 1 & \text{if } s_T = 1 \end{cases}$$

$$\sigma_T^{s*}(\theta_T) = 0$$

$$\sigma_D^*(s_T) = \alpha_T$$

Proposition 2 presents a pooling equilibrium, in which any Target, regardless of type, sends the same signal. In this case, they pool on sending the unaligned signal. Because I have specified that the Donor is willing to pay for policy alignment, all Targets will receive the unaligned Target's reservation price of α_T . For this strategy profile to constitute an equilibrium, the Donor's prior belief must be that alignment is unlikely, so the expected value of an offer of aid, conditional on receiving an unaligned signal, is high; giving aid that exceeds the true reservation price is a remote possibility. Knowing that the Donor is best off giving aid to all ostensibly unaligned Targets, aligned Targets can always fruitfully lie. This region of the parameter space yields the highest utility for an aligned Target, as lying will always be rewarded. By contrast, an unaligned Target is equally well off in the pooling and separating equilibria. For the Donor, the pooling equilibrium is unfavorable. For a potentially significant portion of Targets (a portion which is increasing in the importance of the Target and decreasing in the reservation price), the Donor gives aid in excess of the true reservation price.

In Appendix A, I formally describe two additional equilibria. In the semi-separating equilibrium (Proposition 3), the aligned Target employs a mixed strategy, signaling dishonestly only some of the time. In response to an unaligned signal, the Donor similarly employs a mixed strategy, sending aid




Figure 2: Equilibrium Plot when Donor is Willing to Pay for Policy Change ($\alpha_D > \alpha_T$)

only some of the time. For this strategy profile to constitute an equilibrium, the Donor's prior must be that policy alignment is relatively likely. Intuitively, an unaligned signal from a Target who is very likely *ex ante* to be aligned is not credible, and the Donor expects that no payment is necessary to secure policy alignment. In Proposition 4, I describe a second separating equilibrium (which I call the priced-out equilibrium) in which both types of Target signal honestly but the Donor gives no aid to either type. This strategy profile is an equilibrium if the unaligned Target's reservation price exceeds the Donor's willingness to pay ($\alpha_T \geq \alpha_D$). In this case, changing the Target's mind is simply too costly, so the Donor opts against offering aid. With no aid to be gained, the Target has no incentive to pay the lying cost and thus signals honestly. Figure 2 displays graphically where in the parameter space each equilibrium occurs, conditional on the Donor being willing to pay for policy change.

Taken together, the model suggests three individually necessary and jointly sufficient conditions for false signaling to be optimal at least some of the time for the Target. First, the Target must

be the aligned type ($\theta_T = 1$). Sending a signal of alignment suggests the reservation price for implementing the Donor’s preferred policy is zero, and thus does not yield aid because the Donor does need to spend resources to achieve its objective. Second, the Donor must be willing to pay for the Target to change policy; stating the same condition differently, the Target’s aid demand must lead them to a reservation price that is less than the maximum price the Donor will pay for policy change.⁶ Third, the costs of lying must be lower than the Target’s reservation price for policy change ($c < \alpha_T$). Otherwise, even if lying “works” in the sense that it yields aid, the disutility incurred from lying exceeds the utility gained from aid. Beyond these conditions, the prior probability of policy alignment p affects the probability that the Target lies; as p increases, the credibility of an unaligned signal ($\Pr[\theta_T = 0 \mid s_T = 0]$) declines and the Donor becomes less likely to pay after an unaligned signal.

One reasonable concern of the model’s applicability to observable political interactions is that the model is a one-shot game, whereas, empirically, political interactions are repeated many times. This is potentially problematic because repeated play might mean that the Target’s policy decisions ultimately reveal their type and make misrepresentation of their preferences ineffective at securing additional aid. This interest in discerning type might lead the Donor to withhold aid even in the face of an unaligned signal. To address this possibility, I extend the model to incorporate dynamic considerations (see Appendix A.2 for full results). The first stage of the model is identical to the baseline; in the second stage, the Donor once again seeks to purchase policy from the Target, with the knowledge that the Target’s type has or has not persisted from the first period. The results suggest that repeated play does not necessarily render misrepresentation of preferences ineffective, although it does narrow the circumstances in which it is reliably profitable. Intuitively, if the Donor offers aid sufficient to purchase policy from the unaligned type, they learn no new information about the Target’s type; that is, they do not know if the Target implemented the Donor’s policy because they

⁶Formally, α_D must be strictly greater than α_T .

share the Donor's policy preference or because of the aid the Donor offered. Because the Donor cares about the Target's policy today and discounts the future, he is not willing to withhold aid in order to learn the Target's true type. The results demonstrate that misrepresentation of preferences is most likely to be persistently effective when the cost of buying policy is low (α_T), when Donor interest is high (α_D), and when the prior probability of alignment is low (p).

Hypotheses

What, then, does the model predict empirically? I focus on two implications. First, when two jointly sufficient conditions are met – high aid demand that the donor is willing to meet and low lying costs – the model predicts that signals of nonalignment should be correlated with higher aid. Second, as the posterior probability of alignment ρ increases, this correlation declines because the signal becomes less believable. Consider the following equation, which is a stylized translation of the implications of the model.

$$r_D = \alpha_T \times (\mathbb{1}_{s_T=0} \times \mathbb{1}_{\alpha_D > \alpha_T} \times \mathbb{1}_{c < \alpha_T}) - \alpha_T \times (\mathbb{1}_{s_T=0} \times \mathbb{1}_{\alpha_D > \alpha_T} \times \mathbb{1}_{c < \alpha_T} \times (1 - \rho))$$

I simplify this equation by choosing a setting where lying costs are low. I also label the parts as they correspond to the hypotheses I introduce below.

$$r_D = \underbrace{\alpha_T \times (\mathbb{1}_{s_T=0} \times \mathbb{1}_{\alpha_D > \alpha_T})}_{\text{Unaligned signals increase aid (H1)}} - \underbrace{\alpha_T \times (\mathbb{1}_{s_T=0} \times \mathbb{1}_{\alpha_D > \alpha_T} \times (1 - \rho))}_{\text{Smaller effect when likelihood of policy alignment is high (H2)}}$$

The full hypotheses are thus

Hypothesis 1 (H1) *Holding foreign policy similarity and lying costs constant, strategic nonresponse elicits higher future aid only for poor states that the donor values.*

Hypothesis 2 (H2) *Holding lying costs constant, strategic nonresponse elicits more aid for poor states that the donor values with low foreign policy similarity with the donor than for those same states with high foreign policy similarity.*

Mapping the Model into Empirics

Empirical setting. I choose the UN General Assembly as the empirical setting and the United States as the donor of interest. First, the US has stated publicly that it pays attention to UNGA votes and that aid decisions are made in part based on UNGA votes. Further, the US has a mechanism by which it designates certain votes as important; on important votes, the US lobbies other countries, and other work has found that foreign aid is more strongly correlated with alignment on those votes (Bailey et al., 2017). I expect that, in general, the US does not necessarily care about UNGA votes *in their own right*; instead, the US uses UNGA votes as a signal of the target's orientation on other issues or towards the US more broadly. Second, UNGA resolutions are not binding and thus voting in a way that is divorced from one's preferences is unlikely to cause severe policy repercussions. This suggests that the cost of lying is sufficiently low that targets may misrepresent their true preferences and that we will thus observe non-separating behavior (in model terms, s_T will not necessarily equal θ_T). Third, votes are observable to two key audiences: the US and the target's winning coalition. It is thus unlikely that the US will misinterpret a signal of nonalignment as a signal of alignment or vice versa, which would create some unmodeled mismatch between how the signal is intended to be perceived and how it is actually perceived. The winning coalition has the opportunity to see its (pro- or anti-US) orientation misrepresented by its country on the international stage and inflict costs if necessary. Depending on the orientation of the winning coalition, then, I expect variation in the magnitude of the cost of lying and the target's reservation price.

Of course, the UNGA is not the only plausible means of signaling for recipient states. For

example, leaders may use public or private statements, diplomatic visits, the signing of relatively flexible international agreements, or a host of other methods to signal their underlying preferences. My model does not solely apply to UNGA voting, but there are pragmatic reasons to focus there, beyond the substantive fit discussed above. First, all countries may vote at the UNGA, allowing observation of countries with different model-relevant characteristics. Some other signaling methods, like diplomatic visits, require the costly participation on the part of another state. Second, data on UNGA voting is centrally available. Focusing on statements, for example, would force the analyst into an unclear decision as to which statements to include. It is not clear, for example, the conditions under which statements to a domestic newspaper are viewed by US policymakers. Third, the ultimate position taken by the recipient state is observable, although any private communication between actors is, of course, unobservable. While statements are public, for example, it need not land on a particular position like a vote must (even if that position is silence).

The UNGA is considered to be of low importance by most scholars of international relations; studies that focus on the UNGA typically justify this focus not by its importance, but its status as the “only forum in which a large number of states meet and vote on a regular basis on issues concerning the international community”, making it a convenient location to attempt to observe and measure state preferences and international cleavages (Voeten, 2000, p. 186). Nonetheless, there is considerable evidence that great powers allocate resources to increase voting alignment (Dreher, Nunnenkamp, and Thiele, 2008). This work usually makes the implicit assumption that donors are interested in maximizing voting alignment, not in maximizing the likelihood of a resolution passing; this appears to be consistent with evidence from diplomatic cables, where reports center on how countries voted, not which resolutions were passed. For example, one diplomatic cable notes that despite the consistent passage of resolutions with anti-Israel language, the US works to reduce the margin by which they pass. Notably, the US does also work to advance, stop, and/or modify

the content of certain resolutions in the committee stage with some success, even on resolutions that would pass a plenary vote by a wide margin, meaning that aid may be sensitive to unobserved pre-plenary activities (Scott, 2006).

Signals of alignment and nonalignment (s_T). I focus on abstentions and absences (which I refer to collectively as nonresponse) as a signal of (non)alignment. Nonresponse is preferable to all votes that differ from the US's for a few reasons. First, they do not signal strong commitment to a different policy; in terms of the model, it suggests that α_T is not necessarily high. If α_T is high, donors do not distribute aid to the target under any circumstances, as their interest in the policy is not sufficiently high ($\alpha_T \geq \alpha_D$). Second, it is less likely to inflict prohibitively high lying costs on the targets if it does not match the type. If the value of c was too high, then targets would never be willing to incur the cost of lying. Third, it is less likely to capture signals that are primarily directed at other audiences, although this is still a possibility. Nonresponse may be used to draw funds from multiple donors at once, while voting in line with Russia may be more about an interaction with Russia than one with the US.

Considerable work differentiates between abstentions and absences. Historically, scholars have argued that abstentions are a purposeful statement of neutrality, while absences are connected to capacity constraints or turmoil at home (Bailey et al., 2017; Voeten, 2013). Recent work challenges this; Morse and Coggins (2024) argue that states use absences strategically to avoid punishment for not supporting powerful partners' initiatives. With respect to the model in this paper, it is not clear that one necessarily maps more cleanly onto the concept of an unaligned signal. Abstentions may more definitively shift the US's beliefs in the direction of nonalignment, but as stronger signals, may signify a more genuine and unaffordable disagreement. Absences may be less obvious signs of nonalignment, but do little to resolve uncertainty about where the state stands and do not signal an immovable policy stance, meaning that giving aid is a best response. Drawing on the strategy

from Morse and Coggins (2024), I identify strategic absences – defined as an absence on a day when the country voted on other resolutions. I create a composite *strategic nonresponse* variable, which groups together abstentions and strategic absences. These represent purposeful attempts by states to signal nonalignment. Later in the paper, I use nonstrategic absences as a placebo test.

I examine the relationship between UNGA votes and aid at the country-year level, but signals of alignment happen at the individual vote level. Taking the country-year as the unit of analysis is a practical consideration, as it is not possible to link individual votes to the commitment of individual aid projects. It is similarly difficult to link UNGA votes to a systematic measure of the US’s policy priorities with a particular target. The assumption of this analysis, then, is that nonresponse on a greater proportion of votes indicates signals of nonalignment across a greater number of US policy priorities. As the number of policy priorities on which there is nonalignment increases, the US must allocate more aid to secure alignment.

Aid transfers (r_D). I use official development assistance from the US and aligned international financial institutions to represent aid in the model (r_D). Specifically, I use the sum of US bilateral aid disbursements, World Bank commitments, and IMF commitments. This measure follows evidence from Dreher, Lang, et al. (2022) that the US uses its influence at those institutions to distribute aid to countries that deviate from US positions at the UNSC. This strategy allows the US to pursue cooperation with less risk of a domestic backlash. I use aid as a percentage of lagged GDP as the outcome, following recent guidance from Chen and Roth (2024), who suggest scaling outcomes with zeros instead of using log-like transformations.⁷ To ease interpretation of the results, I multiply aid as a percentage of GDP by 100.

Prior probability of alignment (p). To approximate the underlying probability of foreign policy alignment (p), I measure the similarity of a given target’s portfolio of alliance commitments to the

⁷Using lagged GDP alleviates concerns that the independent variable may be otherwise correlated with changes in GDP, thereby driving changes in the dependent variable even with aid held constant.

US.⁸ While this conceptualization of foreign policy similarity has been used since at least Bueno de Mesquita (1975), scholars have debated the proper way to measure such similarity (Signorino and Ritter, 1999). In line with guidance from Häge (2011) and Miller (2022), I use Cohen’s κ as my preferred indicator of foreign policy similarity, which corrects for the rarity of alliance ties.⁹ This measure approximates the prior likelihood that the US and a given target hold the same policy preferences. Although alliances are formed strategically, the costliness of forming an alliance makes it less likely that alliances are formed as a deliberate misrepresentation of policy preferences. However, alliances do not cover the same breadth of issues that the UN General Assembly does, meaning that the US cannot simply use alliance commitments to infer a target’s true policy position. Crucially, this variable is conceptually distinct from, although correlated with, a target’s importance to the US. For example, the Bahamas and Haiti exhibit high foreign policy similarity to the US, but are less important partners than some states with lower similarity, like the UK; the Caribbean states’ high similarity is probably driven by geographic considerations and their relatively low propensity to form ties, compared to the UK. Table 1 shows that the measures of similarity are correlated with agreement with the US at the UNGA. This is consistent with Proposition 3, where an increase in the probability of alignment decreases the probability of an aligned type sending an unaligned signal.

Cost of lying (c). I use the critical media score from Varieties of Democracy to capture lying costs. This index measures the number and prominence of media outlets that criticize government. This relates to lying costs in two ways. First, it provides a means for unhappy winning coalition members to publicize their grievances with a signal they dislike. With a critical media, those who are unhappy with the government have an opportunity to express that in a way that may mobilize others against the leader. Second, with a more critical media, it is less difficult for a donor to discern backlash. If

⁸To avoid measuring direct alliance ties, which I use later to approximate α_D , I omit ties between the US and the target in calculation of the similarity scores.

⁹I probe the robustness of my results using Scott’s π , a similar measure with slightly different underlying assumptions. I use Cohen’s κ in the main specifications because it is a better predictor of UNGA voting alignment (Table 1).

the backlash is discernible, the donor may know that the signal was dishonest, undermining the key modeling assumption that the cost is unobservable to the donor. The validation exercise in Table 2 shows that, conditional on country- and year-fixed effects, critical media is uncorrelated with foreign policy similarity and direct alliance ties with the US. This indicates my empirical analysis does not conflate c , p , and α_D .

This is an admittedly imperfect proxy. The cost of lying varies with the issue at hand as well as the nature of the target state's regime. For example, perhaps states with large pro-Palestine or pro-Israel lobbies would be severely punished for dishonest signaling on votes related to the Israel-Palestine conflict, but not on other issues. Democratic leaders may face more trouble with lying costs, as they cannot communicate their intent to all winning coalition members without revealing their dishonesty. My measure of lying costs cannot capture all of this variation, especially that at the vote-level.^{10,11} Critical media should be correlated with lying costs across vote and regime types. Without a media that is willing to be critical to the regime, news about an unpopular vote might not be communicated aside negative perspectives; if the vote would be unpopular on its own, it may not be communicated at all. Alternative approaches that focus more deeply on how certain votes matter in certain countries require stronger assumptions about the preferences that domestic audiences possess over different issues over the study period.

Aid demand and reservation price (α_T). To approximate the target's aid demand (α_T), I use a dummy variable that takes the value 1 if the country is eligible for support from the International Development Association (IDA) of the World Bank in a given year. Countries with this status are below a GNI per capita threshold and are not sufficiently creditworthy to borrow on market terms; there is also an assessment of the extent to which the countries implement policies that promote economic growth and poverty reduction. IDA-eligible countries should have higher aid demand

¹⁰To the extent that issue salience is consistent across units, that variation is also held constant by year fixed effects.

¹¹On regime type, critical media is highly correlated with democracy ($r = 0.84$). I also hold constant democracy, so to the extent that there is independent variation I also hold that constant.

(i.e., lower values of α_T), as they tend to be poorest. Notably, countries that “graduate” from this status are still eligible for World Bank support from the International Bank for Reconstruction and Development (IBRD), but graduation is associated with less favorable aid terms and considerable drops in aid receipts, both from the World Bank and other donors (Dreher and Lohmann, 2015; Galiani et al., 2017). Consistent with the model and existing empirical work, Table 3 shows that IDA eligibility is correlated with aid inflows.

Donor interest and willingness to pay (α_D). To capture the importance of the target to the US (α_D), I use direct alliance ties. This is distinct from the alliance portfolio similarity described above. Although a target may have a similar portfolio of commitments, it may not be sufficiently important to US interests to yield a formal alliance. The presence of an alliance suggests the US is willing to incur potential costs in order to cooperate with that target, consistent with the idea of dedicating aid to a country to yield policy cooperation. Alliances may also endogenously increase US interests by creating economic or security ties that are costly to replace; domestic firms with investments in a partner may punish a leader who undermines their returns, or finding a new partner to cooperate with in a particular region may be costly. To show that this measure does not simply capture the prior, Table 4 shows that, conditional on indirect foreign policy similarity, direct alliance ties are not correlated with voting alignment at the UNGA.¹²

Using these empirical mappings of key model terms, I test for the presence of a relationship between strategic nonresponse at the UNGA and future aid from the US. Recall from the equilibria that there are three necessary conditions for an unaligned signal to yield aid: the cost of a dishonest signal (low c) must be low, the target’s aid demand must be high (low α_T), and the donor’s interest in the target’s policy must be sufficiently high (high α_D). For clarity, I restate the hypotheses,

¹²Note that, unlike in the foreign policy similarity term, I do not exclude alliances that are solely non-aggression pacts. The main rationale for this choice is to increase statistical power; the intersection of US allies beyond non-aggression pacts and IDA-eligible states is quite small. Theoretically, this choice is conservative beyond its power implications; allies with only non-aggression pacts are probably less important to the US on average than those with deeper alliance commitments.

substituting in the empirical mappings of the theoretical concepts.

Hypothesis 1 (H1) *Holding foreign policy similarity and lying costs constant, strategic nonresponse elicits higher future aid only among IDA-eligible US allies.*

Hypothesis 2 (H2) *Holding lying costs constant, strategic nonresponse elicits more aid for IDA-eligible US allies with low foreign policy similarity with the US than for those with high foreign policy similarity.*

Empirical Analysis

I use a selection on observables approach. Using OLS, I estimate the following equation to test H1

$$Aid_{i,t+1} = \beta \cdot NonresponsePct_{it} \times IDA_{it} \times USAlly_{it} + \gamma \cdot \mathbf{x}_{i,t-1} + Y_t + C_i + \varepsilon_{it}$$

where β is the parameter of interest and the percentage of votes in which target i abstained or was strategically absent in a given year is the explanatory variable of interest.^{13, 14} I use aid as a percentage of GDP for target unit i in year $t + 1$ to reduce the likelihood of reverse causality. $\mathbf{x}$ is a vector of pre-treatment covariates and Y and C are time- and target country-fixed effects, respectively. Standard errors are clustered at the target country level. To test H2, I include an interaction with US foreign policy similarity.

My analysis covers the years 1992 to 2021. For alliance ties, I use data collected by Leeds et al. (2002). I omit alliances in the data that are solely defined by non-aggression pacts. These alliances have a weaker correlation with UNGA voting alignment, making them a less fitting approximation of the US's prior. Figure 3 shows the distribution of strategic nonresponse percentages broken down by alliance status. Descriptively, the figure is consistent with the model. If US allies are typically

¹³Note that all of the constituent terms of the interaction are included individually, as are the interactions between the individual terms.

¹⁴I exclude votes on which the US also nonresponds.




Figure 3: Distribution of Strategic Nonresponse by Alliance Status. The distribution in yellow represents strategic nonresponse by non-US allies, while the distribution in blue represent strategic nonresponse by US allies.

more important to the US's high goals (high α_D relative to reservation price α_T), they are more likely to profitably misrepresent their type.

I include several covariates to ensure my estimates match the quantities of interest from my hypotheses and reduce the risk of omitted variable bias. In all specifications, I include the Cohen's κ measure of foreign policy similarity and the clean elections index from V-DEM, holding constant the parameters p and c , respectively. Political and economic contexts shape the interests of countries and thus their voting behavior at the UN; I control for democracy, GDP per capita, and population as they also influence aid flows (Alesina and Dollar, 2000; Brazys and Panke, 2017). Membership on the UN Security Council has been shown to shape the allocation and terms of foreign aid, as members vote on high-stakes resolutions that matter more to major powers than most UNGA resolutions (e.g. Berlin et al., 2022; Jud, 2024); it may also create incentives for states to behave differently in the UNGA. I include the lagged value of my dependent variable, which may reflect past efforts at

policy persuasion via aid and affect the marginal cost of giving future aid. Finally, low state capacity may increase the rate of absences via unpredictable government turnover and staff shortages while affecting aid flows in an indeterminate way, as giving is more risky, but may have higher returns (Voeten, 2013).

Results

Figure 4 displays the baseline results.¹⁵ Models 1 and 3 take the composite strategic nonresponse term as the independent variable, while Models 2 and 4 take the individual strategic absence and abstention terms as separate variables. Models 1 and 2 include only the model-based covariates, while Models 3 and 4 include all covariates described in the prior section. The most informative results are the marginal effects. Consistent with H1, the estimated marginal effect of strategic nonresponse is positive and significant at the 95% confidence level in the models where abstentions and absences are measured together. The estimated marginal effects suggests that aid increases by about 2.6 percentage points of GDP as a recipient goes from never to always nonresponding. More realistically, a one standard deviation increase (0.10, or about 7-8 more nonaligned votes) in UNGA nonresponse by an IDA-eligible US ally is correlated with an increase of 0.26 percentage points of GDP worth of aid, or an increase of about \$71 million.¹⁶ For comparison, the sample mean of the outcome variable is 2.5 percent of GDP coming from summed aid. When strategic nonresponse is decomposed in Models 2 and 4, the estimates of the marginal effects are null but directionally consistent with H1. The point estimates suggest that abstentions yield more aid on average than absences do. As hypothesized, Figure 4 shows that these results do not appear for other values of IDA eligibility and US alliance status, for which the marginal effects estimates are substantively smaller (and sometimes negative) and always indistinguishable from zero; substantively speaking,

¹⁵Results in table form in the appendix Table 5.

¹⁶To give a more representative amount, the \$71 million figure excludes India, a massive outlier in terms of its economic size, from the mean GDP calculation; if we include it the predicted increase is about \$99 million. Recall that this sums US bilateral and World Bank and IMF multilateral flows.

signaling nonalignment does not yield additional aid for states that are not IDA-eligible US allies.

What countries drive these results? Figure 5 displays country-years for which the theory predicts positive effects of strategic nonresponse, including the number of years in which a country is an IDA-eligible US ally (which I'll refer to as "eligible signalers" for brevity). Eligible signalers are concentrated mostly in Asia and the Americas. This is consistent with the low incidence of IDA eligibility in Europe and low US strategic valuation of Africa. The plot also includes the country's mean foreign policy similarity score, which is important for knowing the support of the test of H2 below. There is considerable variation among the treated group, although it does not span the full range of values, displayed in the legend.

Figure 6 displays the results of the test of H2.¹⁷ Again, Models 1 and 3 use the strategic nonresponse variable, and 2 and 4 break out nonresponse into abstentions and strategic absences. Models 1 and 2 include only model-based covariates, while Models 3 and 4 include other covariates. Consistent with H2, the estimated marginal effect of strategic nonresponse increases as foreign policy similarity decreases. Intuitively, low similarity targets much more believably signal that they are unaligned with the donor, but the same signal from a high similarity target does not convince the donor. Similar to the results in the test of H1, the marginal effects are positive and significant only in the analyses with the composite nonresponse independent variable (1 and 3) and the point estimates for abstentions are larger than for absences, but fail to reach statistical significance (Models 2 and 4). The results suggest that a one standard deviation increase in nonresponse by an IDA-eligible US ally with the minimum level of foreign policy similarity would see an aid increase of about \$150 million.

These results are generally consistent with the implications of the model. When states strategically nonrespond, they receive more aid in the next year. Of course, this analysis does not reveal whether such nonresponses were honest or dishonest. However, if such nonresponses are profitable for certain states, this suggests that UNGA voting is an opportunity for states to "play hard to get", or

¹⁷Table 6 display the results in table form.




Figure 4: Marginal Effect of Nonresponse on Summed Aid. Marginal effects are computed using the results from Table 5. The top row displays results with only model-based covariates (i.e., Models 1 and 2), while the bottom row displays results including other covariates (i.e., Models 3 and 4). Error bars represent 95 percent confidence intervals. The rightmost blue dot represents the main estimate of interest.




Figure 5: Eligible Signaler Observations by Country. Countries are represented by the ISO alpha-3 codes. A country-year is an eligible signaler when they are an IDA-eligible US ally.




Figure 6: Marginal Effects of Nonresponse with Summed Aid with Variation in Foreign Policy Similarity. Marginal effects are computed using the results from Table 6. The top row displays results with only model-based covariates (i.e., Models 1 and 2), while the bottom row displays results including other covariates (i.e., Models 3 and 4). To isolate the relationship of interest, the IDA and US ally variables are held constant at 1 in all plots. Shaded regions represent 95 percent confidence intervals.

exaggerate the distance between the donor's policy and the target's own ideal point.

Alternative Explanation: Vote Buying

Endogeneity and reverse causality pose serious threats to inference in this setting. In particular, it may be the case that the US lobbies countries to abstain and then rewards them with aid. In the appendix, I present evidence that this is not the case. Specifically, strategic nonresponse is not correlated with increases in past aid, which would reflect earlier US attempts to purchase the recipient's future nonresponse (Figure 13). This helps to rule out vote buying because relatively static characteristics like size, political distance, and poverty make states appealing vote buying targets, so a vote buying story would expect correlation between current nonresponse and past aid. In addition, this theory would expect that high nonresponders possess systematically different characteristics that make them good targets for vote buying, like small, aid-dependent economies. To explore this possibility, I categorize observations as above or below the median level of nonresponse and construct a balance table of key variables for all units and then for eligible signalers only (Tables 7 and 8). Two relevant descriptive patterns emerge. High nonresponders are not obvious targets for vote buying (i.e., smaller and poorer). They do tend to be poorer on a per capita basis, but are significantly larger with similar sized economies. They also are not more likely to be UNSC members, have lower state capacity (i.e., may not show up to vote as reliably), and are not more aid dependent. Instead, the descriptive statistics match my theoretical expectations: high nonresponders are more likely to be poor US allies with less media freedom. When looking only at eligible signalers, the only statistically significant differences are that high nonresponders have less similar alliance portfolios but vote more in line with the US at the UNGA, relative to low nonresponders. In the appendix, I probe the robustness of my results to alternative specifications by using Scott's π as the approximation of p (Figure 9), disaggregated by donor (Figure 10), and using longer leads on the outcome to potentially better capture aid budgeting dynamics (Figures 11 and 12). The key results

from the main text are unchanged.

Placebo Test: Nonstrategic Absences

As discussed above, scholars have long tied absence on UNGA votes to capacity or stability concerns (Voeten, 2013). This is distinct from strategic absences, which are purposeful attempts by states to avoid revealing their preferences (Morse and Coggins, 2024). These distinct rationales lend themselves to a placebo test. Donors are presumably more inclined to help politically aligned states in crisis than unaligned states. Similar to my theory, this line of reasoning predicts that nonstrategic absence is positively correlated with future aid. However, the marginal effect of nonstrategic absence should increase with foreign policy similarity. If these divergent conditional marginal effects exist, it lends evidence to my theory that strategic absences are indicators that donors must pay to secure policy alignment. I rerun my main analyses with nonstrategic absence as the independent variable. The results are displayed in Figure 7. Consistent with the capacity-related absences story, allies with highly similar foreign policy see the highest impact of nonstrategic absence on aid.




Figure 7: Marginal Effect of Nonstrategic Absence. In the plots in the right-hand column, IDA eligibility and US alliance status are held constant at 1 to isolate the relationship of interest.

Conclusion

My findings raise several implications for scholars of international relations. Notably, they suggest that low stakes actions like UNGA votes do not map neatly onto a state's preferences. Considerable work uses UNGA voting alignment – or some ideal point derived from it – as an indicator of foreign policy alignment (e.g. Tomashevskiy, 2021). This work often suggests that the symbolic nature of UNGA voting makes it an appealing forum for approximating preferences. Comparing UNGA votes to alliance formation, Gartzke (2006) writes that the relatively low-cost nature of UNGA votes means that preferences should be less distorted; similarly, Bailey et al. (2017) suggest that the nonbinding nature of UNGA resolutions means that strategic voting is less likely. My results challenge this assumption; according to my model, the minimal consequences of UNGA voting are the *source* of strategic distortions, as states can misrepresent their preferences in search of aid.

How should scholars measure preferences instead? While this article highlights novel flaws in relying on low-cost actions, there are tradeoffs to other approaches. One such approach is to use international agreements as the basis for estimating preferences (Häge, 2011; Montal et al., 2020). Entering into these agreements is more costly in terms of actual policy impact, but they require the consent of and negotiations between multiple actors to enter into force, meaning that state A may be unable to achieve its preferred outcome due to the actions of state B. UNGA votes still have a number of advantages: they are observable for virtually all country-years, on a common set of issues for all countries, and states entirely control their own votes. In model terms, we would like to only capture votes that are in the separating equilibrium, where states signal honestly. To sustain that equilibrium, the cost of lying must be higher than the target's reservation price. This is not a directly verifiable assumption, but this is more likely to be the case on more salient issues and when media attention is high. Identifying these votes is outside of the scope of this project, but future

research might use data on media coverage of the UN *prior* to a vote, and then focus on the top quartile (or some other threshold) of votes in terms of media coverage.¹⁸ In the interim (and for cases where media coverage data is not available), scholars should check robustness to alternative high-cost measures of state preferences.

To what extent are the results from UNGA votes generalizable to other signaling venues? Signaling could take place in a number of settings, such as diplomatic visits, elite public rhetoric, or purposeful leaks to the media. Like UNGA votes, these actions nominally contain information about a state's foreign policy orientation but are sufficiently low cost that they may be exploited strategically. The key insights of this paper for other signaling venues come from the model's scope conditions for profitable misrepresentation: sufficiently low (but positive) costs and donor willingness to pay. For signaling to be effective, it must be costly. It is plausible that leaders could face backlash from their winning coalition for these kinds of actions, although this may not always be the case. In the US, for example, President Trump was criticized by members of his own party for a press conference with President Vladimir Putin, who interpreted his rhetoric as evidence that Trump might be too soft on an adversary (Fox, 2018). Another concern is the interpretation of a signal. UNGA votes are relatively unambiguous, at least in terms of their content. Other venues may not be so straightforward. For example, a leak may report that a president is considering an unpopular foreign policy action. An outside actor cannot discern if the leak is an attempt by the leader to gauge the popularity of a genuinely considered option, to signal their preferences, or the result of internal politics as one set of staffers attempts to undermine another. On this point, the particular signaling state matters. Donors may be more inclined to act on unclear signals from states that are important to their policy goals.

Generally, this work points to a need for further research on strategic decision-making by

¹⁸It is key to focus on *prior* media coverage to ensure that the measurement is based on *expected* costs of voting a certain way.

aid-receiving countries. While donor-side decision-making is undoubtedly important, targets can manipulate the information environment that donors evaluate.

References

- Alesina, A., & Dollar, D. (2000). Who gives foreign aid to whom and why? [Publisher: Springer]. *Journal of Economic Growth*, 5(1), 33–63.
- Allison, S. (2023). Russia is losing the superpower scramble for africa. *The Continent*, (113), 14–16.
- Bailey, M. A., Strezhnev, A., & Voeten, E. (2017). Estimating dynamic state preferences from united nations voting data. *Journal of Conflict Resolution*, 61(2), 430–456.
- Bayer, P., Marcoux, C., & Urpelainen, J. (2015). When international organizations bargain: Evidence from the global environment facility [Publisher: SAGE Publications Inc]. *Journal of Conflict Resolution*, 59(6), 1074–1100.
- Berlin, M. P., Desai, R. M., & Olofsgård, A. (2022). Trading favors? UN security council membership and subnational favoritism in aid recipients. *The Review of International Organizations*.
- Brazys, S., & Panke, D. (2017). Why do states change positions in the united nations general assembly? [Publisher: SAGE Publications Ltd]. *International Political Science Review*, 38(1), 70–84.
- Brown, K. J. (2023). Why hide? africa’s unreported debt to china. *The Review of International Organizations*.
- Bueno de Mesquita, B. (1975). Measuring systemic polarity. *Journal of Conflict Resolution*, 19(2), 187–216.
- Bueno de Mesquita, B., & Smith, A. (2016). Competition and collaboration in aid-for-policy deals. *International Studies Quarterly*, 60(3), 413–426.
- Callander, S., & Wilkie, S. (2007). Lies, damned lies, and political campaigns. *Games and Economic Behavior*, 60(2), 262–286.

- Chen, J., & Roth, J. (2024). Logs with zeros? some problems and solutions*. *The Quarterly Journal of Economics*, 139(2), 891–936.
- Crisman-Cox, C., & Gibilisco, M. (2021). Estimating signaling games in international relations: Problems and solutions. *Political Science Research and Methods*, 9(3), 565–582.
- Cruz, C., & Schneider, C. J. (2017). Foreign aid and undeserved credit claiming [eprint: <https://onlinelibrary.wiley.com/doi/10.1111/ajps.12345>]. *American Journal of Political Science*, 61(2), 396–408.
- Curtice, T., & Reinhardt, E. (2023). The politics of human rights trade sanctions: Evidence from the african growth and opportunity act. *European Journal of International Relations*.
- Dang, T. A., & Stone, R. W. (2021). Multinational banks and IMF conditionality. *International Studies Quarterly*, 65(2), 375–386.
- Dixon, W. J. (1981). The emerging image of u.n. politics [Publisher: Cambridge University Press]. *World Politics*, 34(1), 47–61.
- Dreher, A., Lang, V., Rosendorff, B. P., & Vreeland, J. R. (2022). Bilateral or multilateral? international financial flows and the dirty-work hypothesis. *The Journal of Politics*, 84(4), 1932–1946.
- Dreher, A., & Lohmann, S. (2015). Aid and growth at the regional level [Publisher: Oxford University Press]. *Oxford Review of Economic Policy*, 31(3), 420–446.
- Dreher, A., Nunnenkamp, P., & Thiele, R. (2008). Does US aid buy UN general assembly votes? a disaggregated analysis [Publisher: Springer]. *Public Choice*, 136(1), 139–164.
- Eligon, J. (2023). U.s. lawmakers ask white house to punish south africa for supporting russia. *The New York Times*.
- Fearon, J. D. (1995). Rationalist explanations for war [Publisher: [MIT Press, University of Wisconsin Press, Cambridge University Press, International Organization Foundation]]. *International Organization*, 49(3), 379–414.

- Ferry, L. L., & Zeitz, A. O. (2024). The power of having powerful friends: Evidence from a new dataset of IMF negotiating missions, 1985-2020. *The Review of International Organizations*.
- Fox, L. (2018). Top republicans in congress break with trump over putin comments. *CNN*.
- Galiani, S., Knack, S., Xu, L. C., & Zou, B. (2017). The effect of aid on growth: Evidence from a quasi-experiment. *Journal of Economic Growth*, 22(1), 1–33.
- Gartzke, E. (2006). The affinity of nations index, 1946-2002.
- Gramer, R. (2023). *Inside the desperate diplomatic efforts to salvage u.s.-south africa ties* [Foreign policy]. Retrieved October 12, 2023, from <https://foreignpolicy.com/2023/07/19/south-africa-russia-putin-ukraine-war-global-south-agoa-biden-diplomacy/>
- Häge, F. M. (2011). Choice or circumstance? adjusting measures of foreign policy similarity for chance agreement. *Political Analysis*, 19(3), 287–305.
- Horning, N. R. (2008). Strong support for weak performance: Donor competition in madagascar. *African Affairs*, 107(428), 405–431.
- Jud, S. (2024). Diffusing risk: Bureaucratic agency, UN security council horse-trading, and the role of co-financing. *International Studies Quarterly*, 68(4), sqae140.
- Leeds, B., Ritter, J., Mitchell, S., & Long, A. (2002). Alliance treaty obligations and provisions, 1815-1944 [Publisher: Routledge .eprint: <https://doi.org/10.1080/03050620213653>]. *International Interactions*, 28(3), 237–260.
- Makhanya, M. (2023). In russia’s sweet nothings, south africa sees an invitation to a soviet reunion. *The Continent*, (113), 11–13.
- Mikulaschek, C. (2021). *The power of the weak: How informal power-sharing shapes the work of the UN security council*.
- Miller, S. V. (2022). {Peacesciencer} : An r package for quantitative peace science research. *Conflict Management and Peace Science*, 39(6), 755–779.

- Montal, F., Potz-Nielsen, C., & Sumner, J. L. (2020). What states want: Estimating ideal points from international investment treaty content [Publisher: SAGE Publications Ltd]. *Journal of Peace Research*, 57(6), 679–691.
- Morrison, K. M. (2009). Oil, nontax revenue, and the redistributive foundations of regime stability. *International Organization*, 63(1), 107–138.
- Morrow, J. D. (1994). *Game theory for political scientists*. Princeton University Press.
- Morse, J. C., & Coggins, B. (2024). Your silence speaks volumes: Weak states and strategic absence in the UN general assembly. *The Review of International Organizations*.
- Plaisted, J. (2009). *64th UN general assembly: East asian and pacific islands' performance*.
- Ramsay, K. W. (2004). Politics at the water's edge: Crisis bargaining and electoral competition [Publisher: SAGE Publications Inc]. *Journal of Conflict Resolution*, 48(4), 459–486.
- Scott, G. (2006). *Unga: Report on african delegations at the 61st ga*.
- Scott, J., & Carter, R. (2019). Distributing dollars for democracy: Changing foreign policy contexts and the shifting determinants of US democracy aid, 1975–2010. *Journal of International Relations and Development*, 22(3), 640–675.
- Signorino, C. S., & Ritter, J. M. (1999). Tau-b or not tau-b: Measuring the similarity of foreign policy positions. *International Studies Quarterly*, 43(1), 115–144.
- Snidal, D., Hale, T., Jones, E., Mertens, C., & Milewicz, K. (2024). The power of the “weak” and international organizations. *The Review of International Organizations*.
- Spiegel, P. (2014). How the euro was saved - FT series. *Financial Times*.
- Steinwand, M. C. (2015). Foreign aid and political stability. *Conflict Management and Peace Science*, 32(4), 395–424.
- Swedlund, H. J. (2017). *The development dance: How donors and recipients negotiate the delivery of foreign aid*. Cornell University Press.

- Tomashevskiy, A. (2021). Economic statecraft by other means: The use and abuse of anti-bribery prosecution. *International Studies Quarterly*, 65(2), 387–400.
- Voeten, E. (2000). Clashes in the assembly [Publisher: [MIT Press, University of Wisconsin Press, Cambridge University Press, International Organization Foundation]]. *International Organization*, 54(2), 185–215.
- Voeten, E. (2013). Data and analyses of voting in the united nations general assembly. In *Routledge handbook of international organization* (pp. 54–66). Routledge.
- Voeten, E. (2022). Who votes with russia? *Good Authority*.
- Winters, M. S. (2024). The role of recipient governments in determining foreign aid flows. In R. M. Desai, S. Devarajan, & J. L. Tobin (Eds.), *Elgar handbook on aid and development*. Edward Elgar Publishing.
- Wolford, S. (2014). Showing restraint, signaling resolve: Coalitions, cooperation, and crisis bargaining [eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1111/ajps.12049>]. *American Journal of Political Science*, 58(1), 144–156.

A Additional Formal Results

A.1 Baseline Model

Proposition 3 (Semi-Separating Equilibrium) *Suppose that the cost of lying is low ($c < \alpha_T$), the prior probability of the Target being aligned with the Donor is high ($p > 1 - \frac{\alpha_T}{\alpha_D}$), and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). The following strategy-belief profile is a semi-separating equilibrium. An unaligned Target will signal truthfully, while an aligned Target plays a mixed strategy in its signal choice. Conditional on observing an unaligned signal, the Donor offers aid some of the time. Aligned type Targets and unaligned type Targets that receive an offer of aid implement $y_T = 1$, while unaligned type Targets who do not receive an offer of aid implement $y_T = 0$.*

$$\begin{aligned}\beta_D^*(p, s_T) &= \begin{cases} \rho & \text{if } s_T = 0 \\ 1 & \text{if } s_T = 1 \end{cases} \\ \sigma_T^{s*}(\theta_T) &= \begin{cases} 0 & \text{if } \theta_T = 0 \\ \pi & \text{if } \theta_T = 1 \end{cases} \\ \sigma_D^*(s_T) &= \begin{cases} 0 \text{ with probability } 1 - q & \text{if } s_T = 0 \\ \alpha_T \text{ with probability } q & \text{if } s_T = 0 \\ 0 & \text{if } s_T = 1 \end{cases}\end{aligned}$$

where $\rho \equiv \Pr[\theta_T = 0 \mid s_T = 0] = \frac{1-p}{(1-p)+\pi \cdot p}$, $\pi \equiv \Pr[s_T = 0 \mid \theta_T = 1] = \frac{(1-p)(\alpha_D - \alpha_T)}{\alpha_T \cdot p}$, and $q \equiv \Pr[r_D = \alpha_T \mid s_T = 0] = \frac{c}{\alpha_T}$.

Proposition 3 introduces a semi-separating equilibrium. In my case, aligned targets sometimes truthfully signal and sometimes lie. Intuitively, aligned targets want to blend in with the nonaligned types, but the existence of an nonaligned type is sufficiently unlikely that the donor does not always “believe” the signal and reward it with aid. The mixed strategy response allows the donor to capture some truly unaligned targets while wasting resources less often than they would if they paid all targets that signaled that they were nonaligned. Donors are more likely to give aid to possibly nonaligned targets when the cost of lying is higher, but less likely to do so when the requisite amount of aid is higher (α_T).

Proposition 4 (Priced-Out Equilibrium) *Suppose that the Donor is not willing to pay for policy change ($\alpha_D \leq \alpha_T$). There is a separating equilibrium in which all types of the Target signal truthfully and the Donor never gives aid. The Target will implement the policy that matches its type.*

$$\begin{aligned}\beta_D^{****}(p, s_T) &= \begin{cases} 0 & \text{if } s_T = 0 \\ 1 & \text{if } s_T = 1 \end{cases} \\ \sigma_T^{s****}(\theta_T) &= \theta_T \\ \sigma_D^{****}(s_T) &= 0 \text{ for all } s_T\end{aligned}$$

The priced-out equilibrium described in Proposition 4 is a second separating equilibrium. Each type of Target signals truthfully. The truthfulness stems from the prohibitively high reservation price of the unaligned types. Even if an aligned Target could convince the Donor that it is unaligned, the

Donor would not send aid. With no benefit to lying, the Target signals truthfully and implements its most preferred policy.

A.2 Extension: Dynamic Game

Consider an extension game with two time periods, $t \in 1, 2$. In the first period $t = 1$, the setup is identical to the baseline model. The donor now has a second chance to update his beliefs: the target's choice of first period policy y_1 (and how that might have been influenced by the first period aid offer). In $t = 2$, θ_2 is redrawn with probability $p_2 \equiv \Pr[\theta_2 = \theta_1]$ with $p_2 \in (0, 1)$. With probability $1 - p_2$, the type is redrawn according to p_1 , the same commonly known probability from period 1. The realization of p_2 is common knowledge; in other words, both players know if θ will be redrawn. After p_2 is realized, D again makes an offer of aid $r_2 \geq 0$. Observing this offer, T then chooses to implement a policy $y_2 \in \{0, 1\}$. If $y_2 = 1$, T receives r_2 . At the beginning of the game, the period 2 payoffs are discounted by a common discount factor $\delta \in (0, 1)$. I restrict attention to cases where $\alpha_D \geq \alpha_T$ (i.e., the donor is willing to pay for policy change). The new utility functions are thus:

$$\begin{aligned} U_T(s_T, y_1, y_2, r_1, r_2; \theta_1, \theta_2) &= r_1 \cdot y_1 - \mathbb{1}_{y_1 \neq \theta_1} \cdot \alpha_1 - \mathbb{1}_{s_T \neq \theta_1} \cdot c + \delta (r_2 \cdot y_2 - \mathbb{1}_{y_2 \neq \theta_2} \cdot \alpha_T) \\ U_D(r_1, r_2, y_1, y_2) &= \alpha_D \cdot y_1 - r_1 \cdot y_1 + \delta (\alpha_D \cdot y_2 - r_2 \cdot y_2) \end{aligned}$$

The full path of play is thus

1. Nature chooses whether the Target is aligned with the Donor's policy preferences ($\theta_1 \in \{0, 1\}$). Nature chooses an aligned type with probability $p_1 = \Pr[\theta_1 = 1]$. This alignment is observed only by the Target.
2. The Target privately observes its preference θ_1 and selects a signal $s_T \in \{0, 1\}$.
3. The Donor observes the signal (but not θ_T or any cost incurred by the Target) and decides how much aid r_1 (if any) to offer the Target, conditional on the Target's implementation of the Donor's preferred policy $y_1 = 1$.
4. T observes the aid offer r_1 and chooses to implement a policy $y_1 \in \{0, 1\}$. If the Target chooses the Donor's preferred policy $y_1 = 1$, the Donor pays the promised amount of aid r_1 to the Target. The Donor observes the policy change.
5. Period 1 payoffs are realized.
6. Now in the second stage of the game, Nature chooses whether the period two type ($\theta_2 \in \{0, 1\}$) will be equal to θ_1 (with probability $p_2 \in (0, 1)$) or will be redrawn according to p_1 (with probability $1 - p_2$).
7. Both players observe the realization of p_2 . The Donor decides how much aid r_2 (if any) to offer the Target, conditional on the Target's implementation of the Donor's preferred policy $y_2 = 1$.
8. T observes the aid offer r_2 and chooses to implement a policy $y_2 \in \{0, 1\}$. If the Target chooses the Donor's preferred policy $y_2 = 1$, the Donor pays the promised amount of aid r_2 to the Target.
9. Period 2 payoffs are realized.




Figure 8: Game Tree of Second Stage of Dynamic Model. The first stage is identical to the baseline model, displayed in Figure 1.

Figure 8 displays the second stage path of play in game tree form.

Some actions are optimal in all of the equilibria I describe below. First, the logic of Lemma 1 clearly applies to T's period 2 optimal policy choice, as there is no incentive to manipulate beliefs at the end of the game. What is less obvious is whether that logic still applies to the period 1 policy decision. It could be optimal for the Target to deviate from that strategy (i.e., implement an unaligned policy contra to type or a sufficient aid offer to ensure that D offers aid in period 2). This is not the case. The Target discounts future aid and accounts for the possibility that type will be redrawn, so the gains from period 2 aid are never large enough to offset the losses from either rejecting aid or implementing a contra-type unaligned policy in the absence of aid. Therefore, Lemma 1 applies to both the period 1 and period 2 policy decision.

$$y_1^\dagger = y_2^\dagger = \begin{cases} \theta_t & \text{if } r_t < \alpha_T \\ 1 & \text{if } r_t \geq \alpha_T \end{cases}$$

The Donor's optimal offer is also similarly constrained to 0 and α_T as in Lemma 2. The offer strategy in the subgame where type is redrawn is also consistent across equilibria. This makes intuitive sense; the game is effectively reset, so the information gathered from the signal and first period policy is not helpful. The offer strategy in this subgame thus depends only on the prior.

Lemma 3 *In the subgame where θ_2 is redrawn according to p_1 , the Donor's optimal offer strategy is $r_2^\dagger$.*

$$r_2^\dagger(p_1 \equiv \Pr[\theta_2 = 1]) = \begin{cases} \alpha_T & \text{if } p_1 < 1 - \frac{\alpha_T}{\alpha_D} \\ 0 & \text{if } p_1 \geq 1 - \frac{\alpha_T}{\alpha_D} \end{cases}$$

I now describe the features of three equilibria. The path of play depends on the values of different parameters.

Proposition 5 (Separating Equilibrium) *Suppose that the cost of lying is high ($c \geq \alpha_T + \delta \cdot p_2 \cdot \alpha_T$) and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). There is a separating equilibrium in which all types of the Target signal honestly and the Donor pays each type of Target its reservation price. The policies y_1 and y_2 are implemented according to Lemma 1; in equilibrium, both types of*

Target implement $y_1 = y_2 = 1$.

$$\beta_D^\dagger(s_T, y_1) = \begin{cases} s_T & \text{if } t = 1 \\ s_T & \text{if } t = 2 \text{ and } r_1 \geq \alpha_T \\ y_1 & \text{if } t = 2 \text{ and } r_1 < \alpha_T \end{cases}$$

$$\sigma_T^{s\dagger}(\theta_1) = \theta_1$$

$$\sigma_D^\dagger(s_T, y_1) = \begin{cases} \alpha_T & \text{if } t = 1 \text{ and } s_T = 0 \\ 0 & \text{if } t = 1 \text{ and } s_T = 1 \\ \alpha_T & \text{if } t = 2 \text{ and } y_1 = 0 \\ 0 & \text{if } t = 2 \text{ and } y_1 = 1 \end{cases}$$

Proposition 6 (Pooling Equilibrium) Suppose that the cost of lying is low ($c \leq \alpha_T + \delta \cdot p_2 \cdot \alpha_T$), the prior probability of the Target being aligned with the Donor is low ($p_1 \leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T}$), and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). There is a pooling equilibrium in which the Target always signals that they are unaligned with the Donor and the Donor pays all types of the Target the unaligned type's reservation price (α_T). Consistent with Lemma 1, both types of Target implement $y_1 = y_2 = 1$ in equilibrium.

$$\beta^{\dagger\dagger}(s_T, y_1) = \begin{cases} p_1 \leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T} & \text{if } t = 1 \\ p_1 \leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T} & \text{if } t = 2 \text{ and } r_1 \geq \alpha_T \\ y_1 & \text{if } t = 2 \text{ and } r_1 < \alpha_T \end{cases}$$

$$\sigma_T^{s\dagger\dagger} = 0$$

$$\sigma_D^{\dagger\dagger} = \alpha_T$$

Proposition 7 (First Semi-Separating Equilibrium) Suppose that the cost of lying is low ($c \leq \alpha_T + \delta \cdot p_2 \cdot \alpha_T$), the prior probability of the Target being aligned with the Donor is in an intermediate range ($\frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T} < p_1 < 1 - \frac{\alpha_T}{\alpha_D}$), and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). The following strategy-belief profile is a semi-separating equilibrium. An unaligned Target will signal truthfully, while an aligned Target plays a mixed strategy in its signal choice. Conditional on observing an unaligned signal, the Donor offers aid some of the time in the first stage. Aligned type Targets and unaligned type Targets that receive an offer of aid implement $y_1 = 1$, while unaligned type Targets who do not receive an offer of aid implement $y_1 = 0$. In the second stage, the Donor offers aid some of the time to Targets who implemented $y_1 = 0$ and some of the time to Targets who implemented $y_1 = 1$ and received an offer of aid. Again, aligned type Targets and unaligned type Targets that receive an offer of aid implement $y_2 = 1$, while unaligned type Targets who do not receive an offer of aid implement $y_2 = 0$.

$$\beta^{\dagger\dagger\dagger} = \begin{cases} 1 & \text{if } s_T = 1 \\ 0 \text{ with probability } \rho_1 & \text{if } s_T = 0 \text{ and } t = 1 \\ 1 \text{ with probability } (1 - \rho_1) & \text{if } s_T = 0 \text{ and } t = 1 \\ 0 & \text{if } y_1 = 0 \text{ and } t = 2 \\ 1 & \text{if } y_1 = 1, r_1 < \alpha_T, \text{ and } t = 2 \\ 0 \text{ with probability } \rho_2 & \text{if } y_1 = 1, r_1 \geq \alpha_T, \text{ and } t = 2 \\ 1 \text{ with probability } (1 - \rho_2) & \text{if } y_1 = 1, r_1 \geq \alpha_T, \text{ and } t = 2 \end{cases}$$

$$\begin{aligned}
s_T^{\dagger\dagger\dagger} &= \begin{cases} 0 & \text{if } \theta_1 = 0 \\ 0 \text{ with probability } \pi & \text{if } \theta_1 = 1 \\ 1 \text{ with probability } (1 - \pi) & \text{if } \theta_1 = 1 \end{cases} \\
r_1^{\dagger\dagger\dagger} &= \begin{cases} 0 & \text{if } s_T = 1 \\ \alpha_T \text{ with probability } q_1 & \text{if } s_T = 0 \\ 0 \text{ with probability } (1 - q_1) & \text{if } s_T = 0 \end{cases} \\
r_2^{\dagger\dagger\dagger} &= \begin{cases} 0 & \text{if } y_1 = 1 \text{ and } r_1 < \alpha_T \\ \alpha_T & \text{if } y_1 = 0 \\ \alpha_T \text{ with probability } q_2 & \text{if } y_1 = 1 \text{ and } r_1 \geq \alpha_T \\ 0 \text{ with probability } (1 - q_2) & \text{if } y_1 = 1 \text{ and } r_1 \geq \alpha_T \end{cases}
\end{aligned}$$

where $\rho_1 \equiv \Pr[\theta_1 = 0 \mid s_T = 0] = \frac{1-p_1}{1-p_1+\pi \cdot p_1}$, $\rho_2 \equiv \Pr[\theta_1 = 0 \mid y_1 = 1 \text{ and } r_1 \geq \alpha_T] = \frac{1-p_1}{1-p_1+\pi \cdot p_1}$, $\pi \equiv \Pr[s_T = 0 \mid \theta_1 = 1] = \frac{\alpha_D - \alpha_T - p_1 \cdot \alpha_D + p_1 \cdot \alpha_T}{p_1 \cdot \alpha_T}$, $q_1 \equiv \Pr[r_1 = \alpha_T \mid s_T = 0] = \frac{c \cdot q_2 - c}{c - \alpha_T - p_2 \cdot q_2 \cdot \alpha_T}$, and $q_2 \equiv \Pr[r_2 = \alpha_T \mid y_1 = 1 \text{ and } r_1 \geq \alpha_T] = \frac{c + c \cdot q_1 - q_1 \cdot \alpha_T}{p_2 \cdot q_1 + c}$.

Proposition 8 (Second Semi-Separating Equilibrium) *Suppose that the cost of lying is low ($c \leq \alpha_T + \delta \cdot p_2 \cdot \alpha_T$), the prior probability of the Target being aligned with the Donor is high ($p_1 > 1 - \frac{\alpha_T}{\alpha_D}$), and the Donor is willing to pay for policy change ($\alpha_D > \alpha_T$). The following strategy-belief profile is the Target plays a mixed strategy in signal choice, the Donor plays a mixed strategy in its first aid offer, and then does not offer aid in the second period if type is still uncertain. An unaligned Target will signal truthfully, while an aligned Target plays a mixed strategy in its signal choice. Conditional on observing an unaligned signal, the Donor offers aid some of the time in the first stage. Aligned type Targets and unaligned type Targets that receive an offer of aid implement $y_1 = 1$, while unaligned type Targets who do not receive an offer of aid implement $y_1 = 0$. In the second stage, the Donor only offers aid to Targets who implemented $y_1 = 0$ in the first stage. Unaligned type Targets that receive an offer of aid in the second stage implement $y_2 = 0$, while all other Targets implement according to their type.*

$$\begin{aligned}
\beta^{\dagger\dagger\dagger\dagger} &= \begin{cases} 1 & \text{if } s_T = 1 \\ 0 \text{ with probability } \rho & \text{if } s_T = 0 \text{ and } t = 1 \\ 1 \text{ with probability } (1 - \rho) & \text{if } s_T = 0 \text{ and } t = 1 \\ 0 & \text{if } y_1 = 0 \text{ and } t = 2 \\ 1 & \text{if } y_1 = 1, r_1 < \alpha_T, \text{ and } t = 2 \\ p_1 > 1 - \frac{\alpha_T}{\alpha_D} & \text{if } y_1 = 1, r_1 \geq \alpha_T, \text{ and } t = 2 \end{cases} \\
s_T^{\dagger\dagger\dagger\dagger} &= \begin{cases} 0 & \text{if } \theta_1 = 0 \\ 0 \text{ with probability } \pi & \text{if } \theta_1 = 1 \\ 1 \text{ with probability } (1 - \pi) & \text{if } \theta_1 = 1 \end{cases} \\
r_1^{\dagger\dagger\dagger\dagger} &= \begin{cases} 0 & \text{if } s_T = 1 \\ \alpha_T \text{ with probability } q & \text{if } s_T = 0 \\ 0 \text{ with probability } (1 - q) & \text{if } s_T = 0 \end{cases}
\end{aligned}$$

$$r_2^{\dagger\dagger\dagger\dagger} = \begin{cases} 0 & \text{if } y_1 = 1 \text{ and } r_1 < \alpha_T \\ \alpha_T & \text{if } y_1 = 0 \\ 0 & \text{if } y_1 = 1 \text{ and } r_1 \geq \alpha_T \end{cases}$$

where $\rho \equiv \Pr[\theta_1 = 0 \mid s_T = 0] = \frac{1-p_1}{1-p_1+\pi \cdot p_1}$, $\pi \equiv \Pr[s_T = 0 \mid \theta_1 = 0] = \frac{(p_1-1)(\alpha_D - \delta \cdot \alpha_D + \alpha_T(\delta \cdot p_2 - 1))}{p_1(\alpha_D - \alpha_T)}$, and $q \equiv \Pr[r_1 = \alpha_T \mid s_T = 0] = \frac{c}{\alpha_T}$.

B Proofs of Formal Results

B.1 Proof of Lemma 1

Consider the following policy implementation strategy y_T^* played by T.

$$y_T^* = \begin{cases} \theta_T & \text{if } r_D < \alpha_T \\ 1 & \text{if } r_D \geq \alpha_T \end{cases}$$

I check for profitable deviations below. At the indifference point, T will take the aid money and implement $y_T = 1$. First, suppose that $r_D < \alpha_T$.

$$\begin{aligned} U_T(y_T = \theta_T \mid r_D < \alpha_T) &> U_T(y_T = 1 \mid r_D < \alpha_T) \\ p \cdot r_D + (1-p)(0) &> p \cdot r_D + (1-p)(r_D - \alpha_T) \\ r_D &< \alpha_T \quad \checkmark \\ U_T(y_T = \theta_T \mid r_D < \alpha_T) &> U_T(y_T = 0 \mid r_D < \alpha_T) \\ p \cdot r_D + (1-p)(0) &> p(-\alpha_T) + (1-p)(0) \\ r_D &> -\alpha_T \quad \checkmark \end{aligned}$$

Now suppose that $r_D \geq \alpha_T$.

$$\begin{aligned} U_T(y_T = 1 \mid r_D \geq \alpha_T) &\geq U_T(y_T = 1 \mid r_D \geq \alpha_T) \\ p \cdot r_D + (1-p)(r_D - \alpha_T) &\geq p \cdot r_D + (1-p)(0) \\ r_D &\geq \alpha_T \quad \checkmark \\ U_T(y_T = 1 \mid r_D \geq \alpha_T) &\geq U_T(y_T = 0 \mid r_D \geq \alpha_T) \\ p \cdot r_D + (1-p)(r_D - \alpha_T) &\geq p(-\alpha_T) + (1-p)(0) \quad \checkmark \end{aligned}$$

There are no profitable deviations from y_T^* . $\square$

B.1.1 Extending Lemma 1 to the Dynamic Game

First, consider the decision for an aligned type who may reject aid and implement an unaligned policy. I assume that this strategy is perfectly effective (i.e., the donor's beliefs remain that the observation of an unaligned policy always means that the donor believes the target to be unaligned).

$$\begin{aligned} U_T(y_1 = 0 \mid \theta_1 = 1, r_2(y_1 = 0) = \alpha_T, r_2(y_1 = 1) = 0, r_1 = 0) &> \\ U_T(y_1 = 1 \mid \theta_1 = 1, r_2(y_1 = 0) = \alpha_T, r_2(y_1 = 1) = 0, r_1 = 0) & \\ -\alpha_T + \delta[p_2 \cdot \alpha_T + (1-p_2)(\alpha_D - \alpha_T)] &> 0 + \delta[p_2 \cdot 0 + (1-p_2)(0)] \\ \delta \cdot p_2 &> 1 \quad \times \end{aligned}$$

This is never the case because both δ and p_2 are in the $(0, 1)$ interval. Now consider an unaligned type who rejects an aid offer to secure aid in the second period.

$$\begin{aligned}
 U_T(y_1 = 0 \mid \theta_1 = 0, r_2(y_1 = 0) = \alpha_T, r_2(y_1 = 1) = 0, r_1 = \alpha_T) &> \\
 U_T(y_1 = 1 \mid \theta_1 = 0, r_2(y_1 = 0) = \alpha_T, r_2(y_1 = 1) = 0, r_1 = \alpha_T) & \\
 0 + \delta[p_2 \cdot (\alpha_T - \alpha_T) + (1 - p_2)(\alpha_D - \alpha_T)] &> 0 + \delta[p_2 \cdot (0) + (1 - p_2)(\alpha_D - \alpha_T)] \\
 0 &> 0 \quad \times
 \end{aligned}$$

There is thus no incentive to deviate from the strategy described in Lemma 1 in the first period of the dynamic game. $\square$

B.2 Proof of Lemma 2

Suppose not; that is, suppose that there exists an offer $r_D = \alpha_T \pm \varepsilon$, where $\varepsilon > 0$. I demonstrate below that $r_D = \alpha_T - \varepsilon$ and $r_D = \alpha_T + \varepsilon$ are strictly dominated strategies.

$$\begin{aligned}
 U_D(r_D = 0 \mid y_T^*) &> U_D(r_D = \alpha_T - \varepsilon \mid y_T^*) \\
 p \cdot \alpha_D + (1 - p)(0) &> p(\alpha_D - (\alpha_T - \varepsilon)) + (1 - p)(0) \\
 \alpha_T &> \varepsilon \quad \checkmark
 \end{aligned}$$

$r_D = \alpha_T - \varepsilon$ is thus strictly dominated by $r_D = 0$.

$$\begin{aligned}
 U_D(r_D = \alpha_T \mid y_T^*) &> U_D(r_D = \alpha_T + \varepsilon \mid y_T^*) \\
 \alpha_D - \alpha_T &> \alpha_D - (\alpha_T + \varepsilon) \\
 \varepsilon &> 0 \quad \checkmark
 \end{aligned}$$

There are thus only two possible offers, $r_D = 0$ and $r_D = \alpha_T$, that are not strictly dominated strategies, so $r_D \in \{0, \alpha_T\}$. $\square$

B.3 Proof of Proposition 1

I solve for the equilibrium of this game given the results from Lemmas 1 and 2. Suppose that $s_T^* = \theta_T$ and

$$r_D^* = \begin{cases} 0 & \text{if } s_T = 1 \\ \alpha_T & \text{if } s_T = 0 \end{cases}$$

I check if either player has profitable deviations, beginning with the donor.

$$\begin{aligned}
 U_D(r_D = \alpha_T \mid \beta_D^*, s_T = 0) &\geq U_D(r_D = 0 \mid \beta_D^*, s_T = 0) \\
 \alpha_D - \alpha_T &\geq 0 \\
 \alpha_D &\geq \alpha_T \quad \checkmark \\
 U_D(r_D = 0 \mid \beta_D^*, s_T = 1) &\geq U_D(r_D = \alpha_T \mid \beta_D^*, s_T = 1) \\
 \alpha_D &\geq \alpha_D - \alpha_T \\
 \alpha_T &\geq 0 \quad \checkmark
 \end{aligned}$$

$$\begin{aligned}
U_T(s_T = 1 \mid \theta_T = 1, \beta_D^*, r_D^*) &\geq U_T(s_T = 0 \mid \theta_T = 1, \beta_D^*, r_D^*) \\
0 &\geq \alpha_T - c \\
c &\geq \alpha_T \\
U_T(s_T = 0 \mid \theta_T = 0, \beta_D^*, r_D^*) &\geq U_T(s_T = 1 \mid \theta_T = 0, \beta_D^*, r_D^*) \\
\alpha_T - \alpha_T &\geq -c \\
0 &\geq -c \quad \checkmark
\end{aligned}$$

T has a profitable deviation iff $c \geq \alpha_T$. This equilibrium therefore exists only if $c \geq \alpha_T$. $\square$

B.4 Proof of Proposition 2

I first derive the prior p for which the expected value of an offer of $r_D = \alpha_T$ without an informative signal is greater than that of $r_D = 0$.

$$\begin{aligned}
U_D(r_D = \alpha_T) &> U_D(r_D = 0) \\
\alpha_D - \alpha_T &> p \cdot \alpha_D + (1 - p)(0) \\
1 - \frac{\alpha_T}{\alpha_D} &> p
\end{aligned}$$

Suppose that $\beta_D^{**}(\theta_T) = p \leq 1 - \frac{\alpha_T}{\alpha_D}$, $p < 1 - \frac{\alpha_T}{\alpha_D}$, $s_T^{**} = 0$, and

$$r_D^{**} = \begin{cases} 0 & \text{if } s_T = 1 \\ \alpha_T & \text{if } s_T = 0 \text{ and } \alpha_D > \alpha_T \end{cases}$$

I first check the donor's incentive to deviate.

$$\begin{aligned}
U_D(r_D = \alpha_T \mid s_T = 0, \beta_D^{**}) &\geq U_D(r_D = \alpha_T \mid s_T = 0, \beta_D^{**}) \\
\alpha_D - \alpha_T &\geq p(\alpha_D) + (1 - p)(0) \\
1 - \frac{\alpha_T}{\alpha_D} &\geq p \quad \checkmark \\
U_D(r_D = \alpha_T \mid s_T = 1, \beta_D^{**}) &\geq U_D(r_D = \alpha_T \mid s_T = 1, \beta_D^{**}) \\
\alpha_D &\geq \alpha_D - \alpha_T \\
\alpha_T &\geq 0 \quad \checkmark
\end{aligned}$$

I now check the target's incentive to deviate.

$$\begin{aligned}
U_T(s_T = 0 \mid \theta_T = 1, r_D^{**}, \beta_D^{**}) &\geq U_T(s_T = 1 \mid \theta_T = 1, r_D^{**}, \beta_D^{**}) \\
\alpha_T - c &\geq 0 \\
\alpha_T &\geq c \\
U_T(s_T = 0 \mid \theta_T = 0, r_D^{**}, \beta_D^{**}) &\geq U_T(s_T = 1 \mid \theta_T = 0, r_D^{**}, \beta_D^{**}) \\
\alpha_T &\geq -c \quad \checkmark
\end{aligned}$$

For this equilibrium to hold, the cost of lying must be sufficiently low ($c \leq \alpha_T$). $\square$

B.5 Proof of Proposition 3

Consider the model when $p > 1 - \frac{\alpha_T}{\alpha_D}$ and $\alpha_D > \alpha_T$. I define $\rho = \Pr[\theta_T = 0 \mid s_T = 0]$, $\pi = \Pr[s_T = 0 \mid \theta_T = 1] \in (0, 1)$ and $q = \Pr[r_D = \alpha_T \mid s_T = 0] \in (0, 1)$. Suppose that D holds the

following beliefs:

$$\beta_D^{***}(\theta_T) = \begin{cases} 1 & \text{if } s_T = 1 \\ 0 \text{ with probability } \rho & \text{if } s_T = 0 \\ 1 \text{ with probability } 1 - \rho & \text{if } s_T = 0. \end{cases}$$

I conjecture the following strategies:

$$s_T^{***} = \begin{cases} 0 & \text{if } \theta_T = 0 \\ 0 \text{ with probability } \pi & \text{if } \theta_T = 1 \\ 1 \text{ with probability } 1 - \pi & \text{if } \theta_T = 1 \end{cases}$$

$$r_D^{***} = \begin{cases} 0 & \text{if } s_T = 1 \\ \alpha_T \text{ with probability } q & \text{if } s_T = 0 \\ 0 \text{ with probability } 1 - q & \text{if } s_T = 0 \end{cases}$$

I first solve for the posterior belief that a target claiming to be unaligned is actually unaligned, using Bayes' rule.

$$\begin{aligned} \Pr[\theta_T = 0 \mid s_T = 0] &= \frac{\Pr[s_T = 0 \mid \theta_T = 0] \cdot \Pr[\theta_T = 0]}{\Pr[s_T = 0 \mid \theta_T = 0] \cdot \Pr[\theta_T = 0] + \Pr[s_T = 0 \mid \theta_T = 1] \cdot \Pr[\theta_T = 1]} \\ &= \frac{1 \cdot (1 - p)}{1 \cdot (1 - p) + \pi \cdot p} \\ \rho &= \frac{1 - p}{(1 - p) + \pi \cdot p} \end{aligned}$$

Using this probability, I solve for the mixing probability by an aligned target ($\theta_T = 1$) necessary to make the donor indifferent between giving $r_D = \alpha_T$ and $r_D = 0$, conditional on observing $s_T = 0$.

$$\begin{aligned} U_D(r_D = \alpha_T \mid s_T = 0) &= U_D(r_D = 0 \mid s_T = 0) \\ \alpha_D - \alpha_T &= \Pr[\theta_T = 0 \mid s_T = 0] \cdot (0) + \Pr[\theta_T = 1 \mid s_T = 0] \cdot (\alpha_D) \\ &= \left(\frac{1 - p + \pi \cdot p}{1 - p + \pi \cdot p} - \frac{(1 - p)}{1 - p + \pi \cdot p} \right) (\alpha_D) \\ \alpha_D - \alpha_T &= \frac{\pi \cdot p \cdot \alpha_D}{1 - p + \pi \cdot p} \\ \pi &= \frac{(1 - p)(\alpha_D - \alpha_T)}{\alpha_T \cdot p} \end{aligned}$$

This probability π is between 0 and 1 as long as $p > 1 - \frac{\alpha_T}{\alpha_D}$, which is the same probability that forms the upper bound of the pooling equilibrium. I now determine the mixing probability q for the donor's strategy that makes the aligned type target indifferent between lying and telling the truth.

$$\begin{aligned} U_T(s_T = 1 \mid \theta_T = 1) &= U_T(s_T = 0 \mid \theta_T = 1) \\ 0 &= q \cdot (\alpha_T - c) + (1 - q)(-c) \\ q &= \frac{c}{\alpha_T} \end{aligned}$$

which is in the interval $(0, 1)$ when $c < \alpha_T$.

Because there are neither pooling nor separating strategies that can be sustained in this region of the parameter space and both actors are set indifferent when mixing with probabilities in the interval $(0, 1)$, this is a PBE (Morrow, 1994). $\square$

B.6 Proof of Proposition 4

When $\alpha_T > \alpha_D$, the minimum offer D must make to convince a T with $\theta_T = 0$ to implement $y_T = 1$ is higher than the additional payoff D will receive from policy alignment.

$$\begin{aligned} U_D(r_D = 0 \mid s_T = 0, \alpha_T > \alpha_D) &\geq U_D(r_D = \alpha_T \mid s_T = 0, \alpha_T > \alpha_D) \\ 0 &\geq \alpha_D - \alpha_T \\ \alpha_T &\geq \alpha_D \end{aligned}$$

D will thus never make an offer of aid, as those states that need to be purchased to achieve policy alignment are too pricey ($r_D^{****} = 0$). Because no aid is being offered, no targets have any incentive to deviate from their type ($s_T^{****} = \theta_T$). $\square$

B.7 Proof of Lemma 3

Consider D's optimal period 2 offer in the subgame where θ_2 is redrawn according to p_1 . Recall that both players know that θ_2 is redrawn according to a commonly known p_1 . Also recall that Lemma 2 establishes that there are only two unique equilibrium offers: $r_2 \in \{0, \alpha_T\}$.

$$r_2^*(p_1 \equiv \Pr[\theta_2 = 1]) = \begin{cases} \alpha_T & \text{if } p_1 < 1 - \frac{\alpha_T}{\alpha_D} \\ 0 & \text{if } p_1 \geq 1 - \frac{\alpha_T}{\alpha_D} \end{cases}$$

$$\begin{aligned} U_D(r_2 = \alpha_T) &\geq U_D(r_2 = 0) \\ \alpha_D - \alpha_T &\geq p_1(\alpha_D) + (1 - p_1)(0) \\ p_1 &\leq 1 - \frac{\alpha_T}{\alpha_D} \end{aligned}$$

D thus offers aid at lower reservation prices (α_T), lower probability of alignment (p_1), and higher levels of interest in the target's policy (α_D). $\square$

B.8 Proof of Proposition 5

First, consider D's second period offer when type is not redrawn ($\theta_1 = \theta_2$).

$$\begin{aligned} U_D(r_2 = 0 \mid y_1 = 1) &\geq U_D(r_2 = \alpha_T \mid y_1 = 1) \\ \alpha_D &\geq \alpha_D - \alpha_T \quad \checkmark \\ U_D(r_2 = \alpha_T \mid y_1 = 0) &\geq U_D(r_2 = 0 \mid y_1 = 0) \\ \alpha_D - \alpha_T &\geq 0 \\ \alpha_D &\geq 0 \quad \checkmark \end{aligned}$$

Next, consider T's first period policy decision. First, I show that there is no incentive to defect in the case where $p_1 \leq 1 - \frac{\alpha_T}{\alpha_D}$ (i.e., where $r_1 = \alpha_T$ when θ_2 is redrawn).

$$\begin{aligned} U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = 0) &\geq U_T(y_1 = 0 \mid \theta_1 = 1, r_1 = 0) \\ 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1(\alpha_T) + (1 - p_1)(0))] &\geq -\alpha_T + \delta [p_2 \alpha_T + (1 - p_2)(p_1 \alpha_T + (1 - p_1)(0))] \\ 1 &\geq \delta \cdot p_2 \quad \checkmark \\ U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 0, r_1 = \alpha_T) \\ 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1(\alpha_T) + (1 - p_1)(0))] &\geq 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \alpha_T + (1 - p_1)(0))] \\ 0 &\geq 0 \quad \checkmark \end{aligned}$$

There is no incentive to defect because δ and p_2 are both probabilities in the interval $(0, 1)$. Similarly, there is no incentive to defect in the case where $p_1 > 1 - \frac{\alpha_T}{\alpha_D}$ (i.e., where $r_1 = 0$ when θ_2 is redrawn).

$$\begin{aligned}
 U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = 0) &\geq U_T(y_1 = 0 \mid \theta_1 = 1, r_1 = 0) \\
 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1(0) + (1 - p_1)(0))] &\geq -\alpha_T + \delta [p_2 \alpha_T + (1 - p_2)(p_1(0) + (1 - p_1)(0))] \\
 \alpha_T &\geq 0 \quad \checkmark \\
 U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 0, r_1 = \alpha_T) \\
 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1(0) + (1 - p_1)(0))] &\geq 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1(0) + (1 - p_1)(0))] \\
 0 &\geq 0 \quad \checkmark
 \end{aligned}$$

I now turn to D's first period offer, beginning with the case where $p_1 \leq 1 - \frac{\alpha_T}{\alpha_D}$ (i.e., where $r_1 = \alpha_T$ when θ_2 is redrawn).

$$\begin{aligned}
 U_D(r_1 = 0 \mid s_T = 1) &\geq U_D(r_1 = \alpha_T \mid s_T = 1) \\
 \alpha_D + \delta [p_2 \cdot \alpha_D + (1 - p_2)(\alpha_D - \alpha_T)] &\geq \alpha_D - \alpha_T + \delta [p_2 \cdot \alpha_D + (1 - p_2)(\alpha_D - \alpha_T)] \\
 \alpha_T &\geq 0 \quad \checkmark \\
 U_D(r_1 = \alpha_T \mid s_T = 0) &\geq U_D(r_1 = 0 \mid s_T = 0) \\
 \alpha_D - \alpha_T + \delta [p_2(\alpha_D - \alpha_T) + (1 - p_2)(\alpha_D - \alpha_T)] &\geq 0 + \delta [p_2(\alpha_D - \alpha_T) + (1 - p_2)(\alpha_D - \alpha_T)] \\
 \alpha_D &\geq \alpha_T \quad \checkmark
 \end{aligned}$$

Recall that I have restricted attention to cases where $\alpha_D \geq \alpha_T$. There is not an incentive to deviate in the case where $p > 1 - \frac{\alpha_T}{\alpha_D}$.

$$\begin{aligned}
 U_D(r_1 = 0 \mid s_T = 1) &\geq U_D(r_1 = \alpha_T \mid s_T = 1) \\
 \alpha_D + \delta [p_2 \cdot \alpha_D + (1 - p_2)(p_1 \cdot \alpha_D + (1 - p_1)(0))] &\geq \\
 \alpha_D - \alpha_T + \delta [p_2 \cdot \alpha_D + (1 - p_2)(p_1 \cdot \alpha_D + (1 - p_1)(0))] &\geq \\
 \alpha_T &\geq 0 \\
 U_D(r_1 = \alpha_T \mid s_T = 0) &\geq U_D(r_1 = 0 \mid s_T = 0) \\
 \alpha_D - \alpha_T + \delta [p_2(\alpha_D - \alpha_T) + (1 - p_2)(p_1 \cdot \alpha_D + (1 - p_1)(0))] &\geq \\
 0 + \delta [p_2(\alpha_D - \alpha_T) + (1 - p_2)(p_1 \cdot \alpha_D + (1 - p_1)(0))] &\geq \\
 \alpha_D &\geq \alpha_T
 \end{aligned}$$

Finally, I examine T's choice of signal, beginning with $p_1 \leq 1 - \frac{\alpha_T}{\alpha_D}$. This reveals a constraint on the equilibrium related to c .

$$\begin{aligned}
 U_T(s_T = 1 \mid \theta_1 = 1) &\geq U_T(s_T = 0 \mid \theta_1 = 1) \\
 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq \alpha_T - c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] \\
 c &\geq \alpha_T + \delta \cdot p_2 \cdot \alpha_T \\
 U_T(s_T = 0 \mid \theta_1 = 0) &\geq U_T(s_T = 1 \mid \theta_1 = 0) \\
 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq -c + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] \\
 c &\geq 0
 \end{aligned}$$

The identical condition exists for the case where $p > 1 - \frac{\alpha_T}{\alpha_D}$.

$$\begin{aligned} U_T(s_T = 1 \mid \theta_1 = 1) &\geq U_T(s_T = 0 \mid \theta_1 = 1) \\ 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot 0 + (1 - p_1)(0))] &\geq \alpha_T - c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(p_1 \cdot 0 + (1 - p_1)(0))] \\ c &\geq \alpha_T + \delta \cdot p_2 \cdot \alpha_T \end{aligned}$$

There is only one additional constraint necessary to sustain this separating equilibrium: $c \leq \delta \cdot p_2 \cdot \alpha_T$.

□

B.9 Proof of Proposition 6

First, I consider D's offer in the second period.

$$\begin{aligned} U_D(r_2 = \alpha_T \mid r_1 = \alpha_T, y_1 = 0) &\geq U_D(r_2 = 0 \mid r_1 = \alpha_T, y_1 = 0) \\ \alpha_D - \alpha_T &\geq p_1(\alpha_D) + (1 - p_1)(0) \\ p &\leq 1 - \frac{\alpha_T}{\alpha_D} \\ U_D(r_2 = 0 \mid r_1 = 0, y_1 = 1) &\geq U_D(r_2 = \alpha_T \mid r_1 = 0, y_1 = 1) \\ \alpha_D &\geq \alpha_D - \alpha_T \\ \alpha_T &\geq 0 \end{aligned}$$

There is an incentive to deviate if $p_1 \leq 1 - \frac{\alpha_T}{\alpha_D}$. This is the same cutpoint for D to offer aid in the scenario where θ_2 is redrawn according to p_1 .

Now I check T's period 1 policy decision, beginning with the case where $p \leq 1 - \frac{\alpha_T}{\alpha_D}$.

$$\begin{aligned} U_T(y_1 = 1 \mid \theta_1 = 0, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 0, r_1 = \alpha_T) \\ 0 + \delta [p_2(0) + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq 0 + \delta [p_2(0) + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] \\ 0 &\geq 0 \quad \checkmark \\ U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 1, r_1 = \alpha_T) \\ \alpha_T - c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq -c + \delta [p_2(0) + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] \\ 1 &\geq \delta \cdot p_2 \quad \checkmark \end{aligned}$$

Because δ and p_2 are both probabilities in the interval $(0, 1)$, there is no incentive for T to deviate at this stage.

Now the case where $p_1 > 1 - \frac{\alpha_T}{\alpha_D}$. There is no incentive to deviate in this case.

$$\begin{aligned} U_T(y_1 = 1 \mid \theta_1 = 0, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 0, r_1 = \alpha_T) \\ 0 + \delta [p_2(0) + (1 - p_2)(0)] &\geq 0 + \delta [p_2(0) + (1 - p_2)(0)] \\ 0 &\geq 0 \quad \checkmark \\ U_T(y_1 = 1 \mid \theta_1 = 1, r_1 = \alpha_T) &\geq U_T(y_1 = 0 \mid \theta_1 = 1, r_1 = \alpha_T) \\ \alpha_T - c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(0)] &\geq -c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(0)] \\ \alpha_T &\geq 0 \quad \checkmark \end{aligned}$$

Next, I consider D's period 1 aid offer decision with $p_1 \leq 1 - \frac{\alpha_T}{\alpha_D}$. There is a new constraint on p_1 at this point.

$$\begin{aligned}
 U_D(r_1 = \alpha_T \mid s_T = 0) &\geq U_D(r_1 = 0 \mid s_T = 0) \\
 \alpha_D - \alpha_T + \delta [\alpha_D - \alpha_T] &\geq p_1 \cdot \alpha_D + (1 - p_1)(0) + \\
 &\quad \delta [p_2(p_1 \cdot \alpha_D) + (1 - p_1)(\alpha_D - \alpha_T) + (1 - p_2)(\alpha_D - \alpha_T)] \\
 p_1 &\leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T} \\
 U_D(r_1 = 0 \mid s_T = 1) &\geq U_D(r_1 = \alpha_T \mid s_T = 1) \\
 \alpha_D + \delta [p_2 \cdot \alpha_D + (1 - p_2)(\alpha_D - \alpha_T)] &\geq \alpha_D - \alpha_T + \delta [p_2 \cdot \alpha_D + (1 - p_2)(\alpha_D - \alpha_T)] \\
 \alpha_T &\geq 0
 \end{aligned}$$

I then compare the two constraints to see which one binds.

$$\begin{aligned}
 \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T} &\geq \frac{\alpha_D - \alpha_T}{\alpha_D} \\
 \alpha_D &\geq \alpha_T
 \end{aligned}$$

Therefore, $p_1 \leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T}$ is the relevant constraint on p_1 in this equilibrium. This also means I don't need to check the case where $p > 1 - \frac{\alpha_T}{\alpha_D}$. Finally, I consider T's signaling decision.

$$\begin{aligned}
 U_T(s_T = 0 \mid \theta_1 = 1) &\geq U_T(s_T = 1 \mid \theta_1 = 1) \\
 \alpha_T - c + \delta [p_2 \cdot \alpha_T + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] \\
 c &\leq \alpha_T + \delta \cdot p_2 \cdot \alpha_T \\
 U_T(s_T = 0 \mid \theta_1 = 0) &\geq U_T(s_T = 1 \mid \theta_1 = 0) \\
 0 + \delta [p_2 \cdot 0 + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0))] &\geq -c + \delta [p_2 \cdot 0 + (1 - p_2) \\
 &\quad (p_1 \cdot \alpha_T + (1 - p_2)(p_1 \cdot \alpha_T + (1 - p_1)(0)))] \\
 c &\geq 0 \quad \checkmark
 \end{aligned}$$

The upper bound on c described here mirrors the lower bound for c in the separating equilibrium. Thus, to sustain this equilibrium it must be the case that $\alpha_D \geq \alpha_T$, $p_1 \leq \frac{\alpha_D - \alpha_T}{\alpha_D + \delta \cdot p_2 \cdot \alpha_T}$, and $c \leq \alpha_T + \delta \cdot p_2 \cdot \alpha_T$.

B.10 Proof of Proposition 7

I first solve for the posterior belief that a target that sends an unaligned signal is genuinely unaligned (ρ_1).

$$\begin{aligned}
 \rho_1 &\equiv \Pr[\theta_1 = 0 \mid s_T = 0] = \frac{\Pr[s_T = 0 \mid \theta_1 = 0] \cdot \Pr[\theta_1 = 0]}{\Pr[s_T = 0]} \\
 &= \frac{\Pr[s_T = 0 \mid \theta_2 = 0] \cdot \Pr[\theta_1 = 0]}{\Pr[\theta_1 = 0] \cdot \Pr[s_T = 0 \mid \theta_1 = 0] + \Pr[\theta_1 = 1] \cdot \Pr[s_T = 0 \mid \theta_1 = 1]} \\
 &= \frac{(1 - p_1)}{(1 - p_1) + p_1 \cdot \pi} \\
 \rho_1 &= \frac{1 - p_1}{1 - p_1 + \pi \cdot p_1}
 \end{aligned}$$

Similarly, I solve for the second period posterior belief that a target that sends an unaligned

signal, receives an offer of α_T , and implements an aligned policy is genuinely unaligned (ρ_2). Recall that this is in the subgame where $\theta_2 = \theta_1$, meaning that the prior probability that $\theta_2 = 1$ is now p_1 .

$$\begin{aligned}
 \rho_2 &\equiv \Pr[\theta_2 = 0 \mid s_T = 0, r_1 \geq \alpha_T, y_1 = 1] = \frac{\Pr[y_1 = 1, r_1 \geq \alpha_T \mid \theta_2 = 0] \cdot \Pr[\theta_2 = 0]}{\Pr[y_1 = 1, r_1 \geq \alpha_T]} \\
 &= \frac{\Pr[y_1 = 1, r_1 \geq \alpha_T \mid \theta_2 = 0] \cdot \Pr[\theta_2 = 0]}{\Pr[\theta_2 = 0] \cdot \Pr[y_1 = 1, r_1 \geq \alpha_T \mid \theta_2 = 0] + \Pr[\theta_2 = 1] \cdot \Pr[y_1 = 1, r_1 \geq \alpha_T \mid \theta_2 = 1]} \\
 &= \frac{q_1 \cdot (1 - p_1)}{q_1 \cdot (1 - p_1) + p_1 \cdot \pi \cdot q_1} \\
 \rho_2 &= \frac{1 - p_1}{1 - p_1 + \pi \cdot p_1}
 \end{aligned}$$

Intuitively, the equivalence between these two probabilities makes sense because, conditional on making an offer of $r_1 \geq \alpha_T$ in period 1, the donor learns no new information about the target's type.

From there I can derive the probability of false signaling by an aligned type that makes D indifferent between offering aid and offering no aid in the second period.

$$\begin{aligned}
 U_D(r_2 = \alpha_T \mid y_1 = 1, r_1 \geq \alpha_T) &= U_D(r_2 = 0 \mid y_1 = 1, r_1 \geq \alpha_T) \\
 \alpha_D - \alpha_T &= \Pr[\theta_1 = 0 \mid y_1 = 1, r_1 \geq \alpha_T] \cdot 0 + \Pr[\theta_1 = 1 \mid y_1 = 1, r_1 \geq \alpha_T] \cdot \alpha_D \\
 \alpha_D - \alpha_T &= \rho_2 \cdot 0 + (1 - \rho_2) \cdot \alpha_D \\
 \pi &= \frac{\alpha_D - \alpha_T - p_1 \cdot \alpha_D + p_1 \cdot \alpha_T}{p_1 \cdot \alpha_T}
 \end{aligned}$$

This value is in the interval $(0, 1)$ if $p_1 < 1 - \frac{\alpha_D}{\alpha_T}$.

Finally, the mixing probabilities for the donor that make the target indifferent between sending an honest and dishonest signal.

$$\begin{aligned}
 U_T(s_T = 1 \mid \theta_1) &= U_T(s_T = 0 \mid \theta_1 = 1) \\
 0 &= q_1 [\alpha_T - c + p_2(q_2(\alpha_T) + (1 - q_2)(0)) + (1 - p_2)(0)] + \\
 &\quad (1 - q_1)(-c + p_2(0) + (1 - p_2)(0)) \\
 q_1 &= \frac{c \cdot q_2 - c}{c - \alpha_T - p_2 \cdot q_2 \cdot \alpha_T} \\
 q_2 &= \frac{c + c \cdot q_1 - q_1 \cdot \alpha_T}{p_2 \cdot q_1 + c}
 \end{aligned}$$

Both of these probabilities are increasing in c , decreasing in α_T , and decreasing in p_2 .

B.11 Proof of Proposition 8

I first solve for the posterior belief that a target that sends an unaligned signal is genuinely unaligned (ρ).

$$\begin{aligned}
 \rho &\equiv \Pr[\theta_1 = 0 \mid s_T = 0] = \frac{\Pr[s_T = 0 \mid \theta_1 = 0] \cdot \Pr[\theta_1 = 0]}{\Pr[s_T = 0]} \\
 &= \frac{\Pr[s_T = 0 \mid \theta_2 = 0] \cdot \Pr[\theta_1 = 0]}{\Pr[\theta_1 = 0] \cdot \Pr[s_T = 0 \mid \theta_1 = 0] + \Pr[\theta_1 = 1] \cdot \Pr[s_T = 0 \mid \theta_1 = 1]} \\
 &= \frac{(1 - p_1)}{(1 - p_1) + p_1 \cdot \pi} \\
 \rho &= \frac{1 - p_1}{1 - p_1 + \pi \cdot p_1}
 \end{aligned}$$

Then, I check the incentives to deviate for D's second period aid offer.

$$\begin{aligned}
 U_D(r_2 = 0 \mid y_1 = 0, r_1 \geq \alpha_T) &\geq U_D(r_2 = \alpha_T \mid y_1 = 0, r_1 \geq \alpha_T) \\
 p_1 \cdot \alpha_D + (1 - p_1)(0) &\geq \alpha_D - \alpha_T \\
 p_1 &\geq 1 - \frac{\alpha_T}{\alpha_D}
 \end{aligned}$$

Note that this aligns with the upper bound on the prior in Proposition ?? . Now I check the optimal mixing probability in the target's signal that makes the Donor indifferent between offering and not offered aid in the first period.

$$\begin{aligned}
 U_D(r_1 = \alpha_T \mid s_T = 0) &= U_D(r_1 = 0 \mid s_T = 0) \\
 \alpha_D - \alpha_T + \delta[p_2(\rho(0) + (1 - \rho)(\alpha_D)) + (1 - p_2)(p_1(\alpha_D) + (1 - p_1)(0))] &= \\
 0 + \delta[p_2(\rho(\alpha_D - \alpha_T) + (1 - \rho)(\alpha_D) + (1 - p_2)(p_1(\alpha_D) + (1 - p_1)(0))] &= \\
 \rho &= \frac{\alpha_D - \alpha_T}{\delta \cdot \alpha_D - \delta \cdot p_2 \cdot \alpha_T} \\
 \frac{1 - p_1}{1 - p_1 + \pi \cdot p_1} &= \frac{\alpha_D - \alpha_T}{\delta \cdot \alpha_D - \delta \cdot p_2 \cdot \alpha_T} \\
 \pi &= \frac{(p_1 - 1)(\alpha_D - \delta \cdot \alpha_D + \alpha_T(\delta \cdot p_2 - 1))}{p_1(\alpha_D - \alpha_T)}
 \end{aligned}$$

Finally, I solve for the mixing probability for the first period offer that makes the aligned target indifferent between signaling honestly and dishonestly.

$$\begin{aligned}
 U_T(s_T = 0 \mid \theta_1 = 1) &= U_T(s_T = 1 \mid \theta_1 = 1) \\
 q(\alpha_T - c + p_2(0) + (1 - p_2)(0)) + (1 - q)(-c + p_2(0) + (1 - p_2)(0)) &= 0 \\
 q &= \frac{c}{\alpha_T} \quad \square
 \end{aligned}$$

C Additional Empirical Results

Table 1: Validation of Foreign Policy Similarity measure

	<i>Dependent variable:</i>	
	Agreement with US at UNGA	
	(1)	(2)
Scott's π similarity	0.116*** (0.015)	
Cohen's κ similarity		0.129*** (0.019)
Country fixed effects	Yes	Yes
Year fixed effects	Yes	Yes
Observations	5111	5111
<i>Note:</i>	*p<0.05; **p<0.01; ***p<0.001	

Table 2: Validation of Lying Cost Measure

	<i>Dependent variable:</i>	
	Foreign policy similarity	
	(1)	(2)
Critical media	-0.005 (0.005)	0.006 (0.021)
Country fixed effects	Yes	Yes
Year fixed effects	Yes	Yes
Observations	4643	4643
<i>Note:</i>	*p<0.05; **p<0.01; ***p<0.001	

Table 3: Validation of Aid Neediness Measure

	<i>Dependent variable:</i>	
	US, WB, IMF aid as pct. of GDP	
	(1)	(2)
IDA dummy	1.353*** (0.349)	1.150** (0.395)
GDP per capita		-0.0002 (0.0001)
Country fixed effects	Yes	Yes
Year fixed effects	Yes	Yes
Observations	3746	3746
<i>Note:</i>	*p<0.05; **p<0.01; ***p<0.001	

Table 4: Validation of Donor Interest Measure

<i>Dependent variable:</i>	
Agreement with US at UNGA	
US ally dummy	0.017 (0.009)
Foreign policy similarity	0.126*** (0.019)
Country fixed effects	Yes
Year fixed effects	Yes
Observations	5111
<i>Note:</i> *p<0.05; **p<0.01; ***p<0.001	

Table 5: Baseline Results.

	<i>Dependent variable:</i>			
	Summed aid as pct. of GDP			
	(1)	(2)	(3)	(4)
US ally	−0.978 (0.512)	−0.945 (0.606)	−0.882 (0.554)	−0.824 (0.625)
IDA	1.079* (0.442)	0.869 (0.508)	0.552 (0.398)	0.405 (0.502)
Nonresponse	0.130 (0.775)		1.384 (1.009)	
Abstentions		−0.904 (1.918)		0.980 (2.058)
Absences		0.506 (1.151)		1.510 (1.334)
FP similarity	−1.791 (1.325)	−1.708 (1.320)	−0.204 (1.225)	−0.180 (1.221)
Critical media	0.298 (0.259)	0.301 (0.259)	0.282 (0.235)	0.283 (0.236)
Lagged summed aid			0.259*** (0.048)	0.259*** (0.048)
State capacity			0.223 (0.230)	0.230 (0.231)
Log(Population)			1.384 (1.080)	1.303 (1.083)
GDP per capita			−0.0001 (0.0001)	−0.0001 (0.0001)
UNSC			−0.222	−0.230

Table 5: Baseline Results.

			(0.134)	(0.135)
Democracy			−1.110 (1.288)	−1.106 (1.289)
IDA × US ally	−0.356 (0.782)	−0.340 (0.927)	−0.439 (0.756)	−0.498 (0.937)
Nonresponse × US ally	−0.440 (0.802)		−2.404* (1.009)	
Nonresponse × IDA	−2.740 (1.500)		−2.208 (1.243)	
Nonresponse × IDA × US ally	5.847* (2.455)		5.836** (1.933)	
Abstentions × US ally		−0.424 (2.212)		−2.814 (2.255)
Abstentions × IDA		−0.148 (2.557)		−0.644 (2.689)
Absences × US ally		0.681 (1.097)		−1.162 (1.406)
Absences × IDA		−3.645* (1.716)		−2.713 (1.601)
Abstentions × IDA × US ally		4.791 (4.455)		6.076 (4.454)
Absences × IDA × US ally		4.879* (2.149)		4.096 (2.469)
<i>Marginal effects</i>				
Nonresponse	2.798*		2.607*	
Abstentions		3.315		3.597
Absences		2.421		1.731
Country fixed effects	Yes	Yes	Yes	Yes
Year fixed effects	Yes	Yes	Yes	Yes
All covariates	No	No	Yes	Yes
Observations	3205	3034	3205	3034
<i>Note:</i>			*p<0.05; **p<0.01; ***p<0.001	

Table 6: Baseline Results with Variation in Prior.

	<i>Dependent variable:</i>			
	Summed aid as pct. of GDP			
	(1)	(2)	(3)	(4)
US ally	−0.873	−0.878	−1.033	−0.841

Table 6: Baseline Results with Variation in Prior.

	(0.645)	(0.570)	(0.593)	(0.521)
IDA	1.020 (1.149)	0.645 (0.494)	0.023 (0.463)	0.275 (0.435)
Nonresponse	1.193 (1.372)		1.539 (1.170)	
Abstentions		0.513 (2.257)		3.287 (2.687)
FP similarity	-5.269 (4.843)	-1.236 (3.068)	-1.961 (4.182)	-0.132 (2.728)
Absences		-0.698 (1.182)		0.470 (0.817)
Critical media	0.174 (0.159)	0.312 (0.260)	0.173 (0.246)	0.300 (0.234)
State capacity			0.947 (1.587)	1.117 (1.209)
Log(Population)			-0.0001* (0.0001)	-0.0001 (0.0001)
GDP per capita			-0.074 (0.104)	-0.232 (0.131)
UNSC			0.044 (0.037)	0.258*** (0.047)
Lagged summed aid			-0.648 (1.315)	-1.150 (1.281)
IDA $\times$ US ally			-0.567 (0.849)	-0.482 (0.915)
Nonresponse $\times$ US ally	0.400 (1.870)		-2.860 (1.864)	
Nonresponse $\times$ IDA	-1.876 (1.669)		-1.507 (1.527)	
Abstention $\times$ US ally		-2.448 (2.672)		-5.282* (2.687)
Abstention $\times$ IDA		-1.024 (3.090)		-2.659 (3.658)
US ally $\times$ FP similarity	6.140 (4.852)	0.619 (3.479)	1.547 (3.590)	0.269 (3.000)
IDA $\times$ FP similarity	2.287 (5.341)	-2.166 (3.689)	0.197 (4.611)	-1.073 (2.747)
Nonresponse $\times$ FP similarity	22.289		13.806	

Table 6: Baseline Results with Variation in Prior.

	(14.191)		(13.110)
Nonresponse $\times$ IDA $\times$ US ally	4.461 (2.759)		7.472** (2.700)
Abstention $\times$ FP similarity		26.740 (20.139)	29.356 (19.345)
Absence $\times$ US ally		1.391 (1.509)	-0.332 (1.047)
Absence $\times$ IDA		-3.401 (2.164)	-2.149 (1.401)
Absence $\times$ FP similarity		-14.923 (16.860)	-11.746 (16.116)
Abstention $\times$ IDA $\times$ US ally		6.942 (5.058)	8.750 (5.394)
IDA $\times$ US ally $\times$ FP similarity	0.024 (6.420)	4.185 (4.502)	3.381 (4.992)
Nonresponse $\times$ US ally $\times$ FP similarity	-23.824 (15.610)		-11.155 (12.542)
Nonresponse $\times$ IDA $\times$ FP similarity	-36.930* (18.266)		-26.574 (18.318)
Nonresponse $\times$ IDA $\times$ US ally $\times$ FP similarity	34.329 (20.716)		18.942 (18.938)
Abstention $\times$ US ally $\times$ FP similarity		-25.509 (21.637)	-29.197 (20.747)
Abstention $\times$ IDA $\times$ FP similarity		-23.424 (27.573)	-27.891 (24.144)
Absence $\times$ IDA $\times$ US ally		5.505** (2.009)	3.982 (2.179)
Absence $\times$ US ally $\times$ FP similarity		15.682 (16.840)	12.206 (15.913)
Absence $\times$ IDA $\times$ FP similarity		0.837 (21.887)	5.724 (19.143)
Abstention $\times$ IDA $\times$ US ally $\times$ FP similarity		19.523 (28.178)	25.200 (24.672)
Absence $\times$ IDA $\times$ US ally $\times$ FP similarity		-7.321 (23.785)	-10.570 (20.378)
<i>Marginal effects at minimum FP similarity</i>			
Nonresponse	4.519*		3.927*
Abstentions		4.879	4.944

Table 6: Baseline Results with Variation in Prior.

Absences		4.714		3.44
Country fixed effects	Yes	Yes	Yes	Yes
Year fixed effects	Yes	Yes	Yes	Yes
All covariates	No	No	Yes	Yes
Observations	3088	2871	3205	3034

Note:

*p<0.05; **p<0.01; ***p<0.001




Figure 9: Marginal Effect of Strategic Nonresponse Using Scott's π for FP Similarity.




Figure 10: Marginal Effect of Strategic Nonresponse with Disaggregated Outcome.




Figure 11: Marginal Effect of Strategic Nonresponse Using Two-Year Lead on Outcome.




Figure 12: Marginal Effect of Strategic Nonresponse Using Three-Year Lead on Outcome.




Figure 13: Placebo Test: Marginal Effect of Strategic Nonresponse on Prior Year's Aid.

	Low Nonresponse (N=1923)		High Nonresponse (N=1277)		Diff. in Means	Std. Error
	Mean	Std. Dev.	Mean	Std. Dev.		
Strategic nonresponse	0.067	0.034	0.209	0.100	0.142***	0.003
Summed aid as pct. of GDP	1.638	3.481	1.690	2.875	0.052	0.113
US ally	0.351	0.477	0.486	0.500	0.134***	0.018
IDA dummy	0.558	0.497	0.609	0.488	0.051**	0.018
US foreign policy similarity	0.053	0.295	0.041	0.244	-0.012	0.010
UNGA agreement with US	0.165	0.085	0.191	0.111	0.026***	0.004
Critical media	0.649	1.225	0.504	1.319	-0.145**	0.046
State capacity	-0.134	1.155	-0.180	1.079	-0.046	0.040
Population (in millions)	40.231	133.486	53.609	201.084	13.378*	6.398
GDP per capita	3770.699	4235.185	3256.528	3435.705	-514.171***	136.276
GDP (in millions constant USD)	135 968.372	648 865.569	155 225.563	786 878.122	19 257.190	26 529.439
UNSC dummy	0.069	0.253	0.059	0.235	-0.010	0.009
Democracy	0.450	0.225	0.419	0.204	-0.030***	0.008

Table 7: Balance Table of Low versus High Nonresponding Units. High nonresponse observations are above or equal to the median value of nonresponse for their year.

	Low Nonresponse (N=204)		High Nonresponse (N=257)			
	Mean	Std. Dev.	Mean	Std. Dev.	Diff. in Means	Std. Error
Strategic nonresponse	0.062	0.031	0.226	0.114	0.164***	0.007
Summed aid as pct. of GDP	2.033	2.435	1.987	1.711	-0.046	0.201
US foreign policy similarity	0.289	0.301	0.063	0.204	-0.226***	0.025
UNGA agreement with US	0.177	0.082	0.227	0.107	0.050***	0.009
Critical media	0.422	1.114	0.328	1.425	-0.094	0.118
State capacity	-0.449	0.819	-0.311	0.835	0.139+	0.077
Population (in millions)	27.797	47.319	40.198	162.128	12.402	10.642
GDP per capita	2043.076	1226.280	1954.430	1008.553	-88.646	106.439
GDP (in millions constant USD)	39 758.641	61 691.206	57 281.326	213 641.565	17 522.685	14 009.070
UNSC dummy	0.029	0.169	0.023	0.151	-0.006	0.015
Democracy	0.434	0.181	0.410	0.160	-0.025	0.016

Table 8: Balance Table of Low versus High Nonresponding Units, Eligible Signalers Only. High nonresponse observations are above or equal to the median value of nonresponse for their year.